

The Swish Concurrency Engine

Bob Burger, editor

© 2018 Beckman Coulter, Inc. Licensed under the MIT License.

Contents

1	Introduction to Swish	6
1.1	Overview	6
1.2	Supervision Tree	7
2	Operating System Interface	8
2.1	Introduction	8
2.2	Theory of Operation	8
2.3	Programming Interface	9
2.3.1	System Functions	9
3	Erlang Embedding	11
3.1	Introduction	11
3.2	Data Structures	11
3.3	Theory of Operation	13
3.4	Programming Interface	15
3.4.1	Process Creation	15
3.4.2	Process Registration	15
3.4.3	Process Termination, Links, and Monitors	16
3.4.4	Messages and Pattern Matching	18
3.4.5	Process Properties	19
3.4.6	Miscellaneous	20
3.4.7	Tuples	22
3.4.8	I/O	23
3.4.9	Queues	31
3.4.10	Hash Tables	31

3.4.11	Error Strings	32
3.4.12	String Utilities	33
4	Generic Server	35
4.1	Introduction	35
4.2	Theory of Operation	35
4.3	Programming Interface	36
4.4	Published Events	38
4.5	Callback Interface	39
5	Event Manager	42
5.1	Introduction	42
5.2	Theory of Operation	42
5.3	Programming Interface	44
6	Gatekeeper	46
6.1	Introduction	46
6.2	Theory of Operation	46
6.3	Programming Interface	48
7	Supervisor	49
7.1	Introduction	49
7.2	Theory of Operation	49
7.3	Design Decisions	52
7.4	Programming Interface	52
7.5	Published Events	54
7.6	Watcher Interface	55
8	Application	57
8.1	Introduction	57
8.2	Theory of Operation	57
8.3	Programming Interface	58
9	Database Interface	59
9.1	Introduction	59

9.2	Theory of Operation	59
9.3	Design Decisions	61
9.4	Programming Interface	62
10	Log Database	67
10.1	Introduction	67
10.2	Theory of Operation	67
10.2.1	Initialization	67
10.2.2	Extensions	68
10.3	Programming Interface	68
10.4	Published Events	70
11	System Statistics	71
11.1	Introduction	71
11.2	Theory of Operation	71
11.3	Programming Interface	71
11.4	Published Events	72
12	HTTP Interface	73
12.1	Introduction	73
12.2	Theory of Operation	73
12.3	Security	74
12.4	Dynamic Pages	74
12.5	Dynamic Page Constructs	74
12.6	Programming Interface	75
12.6.1	JavaScript Object Notation	79
12.7	Published Events	81
13	Command Line Interface	82
13.1	Introduction	82
13.2	Theory of Operation	82
13.3	Programming Interface	84
	Bibliography	88

List of Figures	90
Index	91

Chapter 1

Introduction to Swish

1.1 Overview

The Swish Concurrency Engine is a framework used to write fault-tolerant programs with message-passing concurrency. It uses the Chez Scheme [6] programming language and embeds concepts from the Erlang [8] programming language. Swish also provides a web server following the HTTP protocol [13].

Swish uses message-passing concurrency and fault isolation to provide fault-tolerant software [1, 16]. The software is divided into lightweight processes that communicate via asynchronous message passing but are otherwise isolated from each other. Because processes share no mutable state, one process cannot corrupt the state of another process—a problem that plagues software using shared-state concurrency.

Exceptions are raised when the software detects an error and cannot continue normal processing. If an exception is not caught by the process that raised it, the process is terminated. An error logger records process crashes and other software errors.

There are two mechanisms for detecting process termination, *links* and *monitors*. Processes can be linked together so that when one exits abnormally, the others are killed. A process can monitor other processes and receive process-down messages that include the termination reason.

A single event dispatcher receives events from the various processes and sends them to all attached event handlers. Event handlers filter events based on their needs.

Swish is written in Chez Scheme for two main reasons. First, it provides efficient first-class continuations [4, 18] needed to implement lightweight processes with much less memory and CPU overhead than operating system threads. Second, Chez Scheme provides powerful syntactic abstraction capabilities [7] needed to make the code closely reflect the various aspects of the design. For example, the message-passing system uses syntactic abstraction to specify pattern matching succinctly.

I/O operations are performed asynchronously using C code (see Chapter 2), and they complete via Scheme callback functions. Asynchronous I/O is used so that Swish can run in a single thread without blocking for I/O. The results from asynchronous operations are invoked synchronously by the Scheme code, allowing it to control re-entrancy.




Figure 1.1: Supervision Tree

1.2 Supervision Tree

By default, Swish uses the supervision tree illustrated in Figure 1.1. The top-level supervisor, `main-sup`, is configured one-for-all and no restarts so that a failure of any of its children crashes the program. The `event-mgr` worker is the event manager gen-server (see Chapter 5). The `log-db` worker is a database gen-server (see Chapter 10) that logs all events to the log database. The `event-mgr-sentry` worker is used during shutdown to make sure the event manager stops sending events to `log-db` before `log-db` shuts down. The `statistics` worker is a system statistics gen-server (see Chapter 11) that periodically posts a `<statistics>` event. The `gatekeeper` worker is the gen-server described in Chapter 6.

The `http-sup` supervisor is configured one-for-one with up to 10 restarts every 10 seconds. The `http-listener`, `http-cache`, and `http-connection` processes are described in Chapter 12.

When running as a Windows service, Swish is configured to restart automatically 10 seconds after a failure. The failure is logged to the log file, and an event is logged to the Windows event log.

Chapter 2

Operating System Interface

2.1 Introduction

This chapter describes the operating system interface. Swish is written in Chez Scheme and runs on Linux, macOS, and Windows. It provides asynchronous I/O via libuv [19] and database support via SQLite [20].

2.2 Theory of Operation

The operating system interface is written in C99 [5] as a shared library that is statically linked against the libuv and SQLite libraries and dynamically linked against single-threaded Chez Scheme 9.5. Please refer to Chapter 4 of the *Chez Scheme Version 9 User's Guide* [6] for information on the foreign function interface.

The single-threaded version of Chez Scheme is used because of its simplicity. All Scheme code runs in the main thread. In order to keep this thread responsive, operations that block for more than a millisecond are performed asynchronously.

For each asynchronous function in the operating system interface, a Scheme callback procedure is passed as the last argument. This callback procedure is later returned to Scheme in a list that includes the results of the asynchronous function call.

Scheme object locking and unlocking is handled by the operating system interface because it manages the data structures that contain pointers to Scheme objects.

The operating system interface uses port objects for files, console input, pipes to other processes, and TCP/IP connections. A port object is created by the various open functions, which return a port handle that is used for read, write, and close operations. Once a port is closed, its port object is freed.

For interface functions that can fail, an error pair (*who* . *errno*) is returned, where *who* is a symbol representing the name of the particular function that failed and *errno* is either an error number or, in the case of certain SQLite functions, a pair whose car is the error number and cdr is the English error string.

Section 2.3 describes the programming interface from the C side. The Scheme library (`osi`) provides foreign procedures for each C function using the same name. For functions that may return error pair (*who* . *errno*), the corresponding Scheme procedure *p*, e.g., `osi_read_port`, raises exception `#(osi-error p who errno)`. In addition, the (`osi`) library exports another procedure with the `*` suffix, e.g., `osi_read_port*`, that returns the error pair.

2.3 Programming Interface

Unless otherwise noted, all C strings are encoded in UTF-8.

2.3.1 System Functions

<code>ptr osi_get_argv(void);</code>	function
--------------------------------------	-----------------

The `osi_get_argv` function returns a Scheme vector of strings constructed from the most recent arguments passed to `osi_set_argv`.

<code>size_t osi_get_bytes_used(void);</code>	function
---	-----------------

The `osi_get_bytes_used` function returns the number of bytes used by the C run-time heap. On Linux, it calls the `mallinfo` function. On macOS, it calls the `mstats` function. On Windows, it calls the `_heapwalk` function.

<code>ptr osi_get_callbacks(uint64_t timeout);</code>	function
---	-----------------

The `osi_get_callbacks` function returns a list of callback lists in reverse order of time received. When the list is empty, it blocks up to *timeout* milliseconds before returning. Each callback list has the form (*callback result* ...), where *callback* is the callback procedure passed to the asynchronous function that returned one or more *results*.

<code>const char* osi_get_error_text(int err);</code>	function
---	-----------------

The `osi_get_error_text` function returns the English string for the given error number.

<code>ptr osi_get_hostname(void);</code>	function
--	-----------------

The `osi_get_hostname` function returns the host name from `uv_os_gethostname`.

<code>uint64_t osi_get_hrttime(void);</code>	function
--	-----------------

The `osi_get_hrttime` function returns the current high-resolution real time in nanoseconds from `uv_hrttime`. It is not related to the time of day and is not subject to clock drift.

`uint64_t osi_get_time(void);` **function**

The `osi_get_time` function returns the current clock time in milliseconds in UTC since the UNIX epoch January 1, 1970. On Windows, it calls the `GetSystemTimeAsFileTime` function in `kernel32.dll`. On all other systems, it calls the `clock_gettime` function with `CLOCK_REALTIME`.

`int osi_is_tick_over(void);` **function**

The `osi_is_tick_over` function returns 1 if the current time from `uv_hrttime` is greater than or equal to the threshold set by the most recent call to `osi_set_tick` and 0 otherwise.

`ptr osi_list_uv_handles(void);` **function**

The `osi_list_uv_handles` function calls `uv_walk` and returns a list of pairs (*handle* . *type*), where *handle* is the address of the `uv_handle_t` and *type* is the `uv_handle_type`.

`ptr osi_make_uuid(void);` **function**

The `osi_make_uuid` function returns a new universally unique identifier (UUID) as a bytevector. On Windows, it calls the `UuidCreate` function in `rpcrt4.dll`. On all other systems, it calls the `uuid_generate` function.

`(string->uuid s)` **procedure**
returns: a UUID bytevector

The `string->uuid` procedure returns the bytevector *uuid* for string *s* such that (`uuid->string uuid`) is equivalent to *s*, ignoring case. If *s* is not a string with uppercase or lowercase hexadecimal digits and hyphens as shown in `uuid->string`, exception `#(bad-arg string->uuid s)` is raised.

`(uuid->string uuid)` **procedure**
returns: a string

The `uuid->string` procedure returns the uppercase hexadecimal string representation of *uuid*, `HH3HH2HH1HH0-HH5HH4-HH7HH6-HH8HH9-HH10HH11HH12HH13HH14HH15`, where *HH_i* is the 2-character uppercase hexadecimal representation of the octet at index *i* of bytevector *uuid*. If *uuid* is not a bytevector of length 16, exception `#(bad-arg uuid->string uuid)` is raised.

`void osi_set_argv(int argc, const char *argv[]);` **function**

The `osi_set_argv` function stores the *argv* pointer to a C vector of *argc* strings for use in the `osi_get_argv` function.

`void osi_set_tick(uint64_t nanoseconds);` **function**

The `osi_set_tick` function sets the threshold for `osi_is_tick_over` to be the current time from `uv_hrttime` plus the given number of *nanoseconds*.

Chapter 3

Erlang Embedding

3.1 Introduction

This chapter describes the design of the message-passing concurrency model. It provides a Scheme embedding of a significant subset of the Erlang programming language [1, 2].¹ Tuple and pattern matching macros provide succinct ways of composing and decomposing data structures.

The basic unit of sequential computation is the *process*. Each process has independent state and communicates with other processes by message passing. Because processes share no mutable state, one process cannot corrupt the state of another process—a problem that plagues software using shared-state concurrency. **Concern:** System procedures that mutate data can cause state corruption. **Mitigation:** The code is inspected for use of these procedures.

An uncaught exception in one process does not affect any other process. A process can be monitored for termination, and it can be linked to another process so that, when either process exits, the other one receives an exit signal. Processes are implemented with one-shot continuations [4], and the concurrent system is simulated by the single-threaded program using software timer interrupts. The operating system interface (see Chapter 2) provides asynchronous input/output (I/O) so that processes waiting for I/O do not stop other processes from executing.

For exceptions, we use Erlang’s approach of encoding the information in a machine-readable datum rather than a formatted string. Doing so makes it possible to write code that matches particular exceptions without having to parse strings, and the exception is human language independent.

The rest of this chapter is organized as follows. Section 3.2 introduces the main data structures, Section 3.3 describes how the concurrency model works, and Section 3.4 gives the programming interface.

3.2 Data Structures

q Queues are used in several key places: the inbox of messages for each process, the list of processes ready to run, and the list of sleeping processes. A *queue* is a doubly-linked list with a sentinel value,

¹Tuples, denoted by $\{e_1, \dots, e_n\}$ in Erlang, are implemented as vectors: $\#(e_1 \dots e_n)$. Similarly records, defined as syntactic sugar over tuples in Erlang, are implemented as syntactic sugar over vectors.

the queue's identity. Both the sentinel value and the elements of the queue are instances of `q`, a Scheme record type with mutable `prev` and `next` fields. This representation enables constant-time insertion and deletion operations.

msg When a *message* is sent to a process, its contents are wrapped in an instance of `msg`, a Scheme record type that extends `q` with an immutable `contents` field. This `msg` is inserted into the process's inbox and removed when the process receives it.

pcb A *process* is an instance of `pcb`, a Scheme record type that extends `q` with an immutable `id` field, the process's unique positive exact integer, an immutable `create-time` field, the process's create time from `erlang:now`, and the following mutable fields:

- **name**: registered name or `#f`
- **cont**: one-shot continuation if live and not currently running or `#f` otherwise
- **sic**: system interrupt count
- **winders**: list of winders if live and not currently running or `()` otherwise
- **exception-state**: exception state if live and not currently running, exit reason if dead, or `#f` if currently running
- **inbox**: queue of `msg` if live or `#f` if dead
- **precedence**: wake time if sleeping or 0 if ready to run
- **flags**: fixnum with bit 0 set when sleeping and bit 1 set when the process traps exits
- **links**: list of linked processes
- **monitors**: list of monitors
- **src**: source location `#(at char-offset filename)` when available if waiting in a `receive` macro or `#f` otherwise

mon A *monitor* is an instance of `mon`, a Scheme record type with two immutable fields, `origin` and `target`, each of which is a process.

osi-port An *osi-port* is an instance of `osi-port`, a Scheme record type with an immutable `name` field and a mutable `handle` field that wraps an operating system interface port. The `handle` field is set to `#f` when the `osi-port` is closed.

directory-watcher A *directory watcher* is an instance of `directory-watcher`, a Scheme record type with a mutable `handle` field and an immutable `path` field. The `handle` field is set to `#f` when the directory watcher is closed.

listener A *TCP listener* is an instance of **listener**, a Scheme record type with immutable **address**, **port-number**, and **create-time** fields and a mutable **handle** field. The **handle** field is set to **#f** when the listener is closed.

3.3 Theory of Operation

The system uses a *scheduler* to execute one process at a time. Each process holds its own system interrupt count (updated by **enable-interrupts** and **disable-interrupts**), list of winders (maintained by **dynamic-wind** and the system primitive **\$current-winders**), and exception state (maintained by **current-exception-state**). The scheduler captures the one-shot continuation for a process with an empty list of winders so that, when it invokes the continuation of another process, it does not run any winders. **Concern:** Using a system procedure that relies on the global winders list may lead to incorrect behavior. **Mitigation:** System procedures that rely on the global winders list are called from only one process at a time using the *gatekeeper* described in Chapter 6. The gatekeeper hooks the **\$cp0**, **\$np-compile**, **pretty-print**, and **sc-expand** system primitives.

Spawning a new process is not as simple as capturing a one-shot continuation and creating a **pcb** record, because the continuation's stack link [18] would be the continuation of the caller, and its list of winders would be the caller's. Thus, the scheduler remembers the current list of winders and then sets it to the empty list before capturing a one-shot continuation. This return continuation is stored in a mutable variable so that it is not closed over by the new process. Next, a full continuation is captured to create the initial exception state that will terminate the new process when an uncaught exception is raised. So that this full continuation does not refer to the caller's continuation, the current stack link is set to the null continuation before capturing it. After capturing the full continuation, a one-shot continuation for the new process is captured and returned to the caller via the return continuation.

Each process runs until it waits in a **receive** macro, is preempted by the **timer-interrupt-handler**, or exits. The operating system interface (see Chapter 2) provides asynchronous I/O operations so that the scheduler can execute other processes while the system is performing I/O. The timer interrupt handler runs every 1000 procedure calls.² The scheduler uses **osi::SetTick** and **osi::IsTickOver** to determine when the time quantum for a process has elapsed.

When process *p* exits with reason *r*, the message **#(DOWN *m* *p* *r*)** is sent to each of its monitor *m*'s **origin** processes. The message **#(EXIT *p* *r*)** is sent to each linked process that traps exits. If *r* is not **normal**, each linked process that does not trap exits is killed with reason *r*.

A process can be registered with a global name, a symbol. This name can be used instead of the process record itself to send it messages. A global *registrar* maintains an eq-hashtable mapping names to processes. The reverse mapping is maintained in the **pcb** record through the **name** field.

There are two system processes: the *event-loop* and the *finalizer*.

The event-loop process calls **osi::GetCompletionPacket** to retrieve completion packets from the operating system interface. It executes each callback with interrupts disabled. Event-loop callbacks are designed to execute quickly without failing or causing new completion packets to be enqueued. Typical callbacks register objects that wrap operating system interface handles with

²1000 was chosen because Chez Scheme performs its internal interrupt checks every 1000 ticks.

a guardian and send messages to a process. If the event-loop process exits with reason *r*, the system logs the event `#(event-loop-process-terminated r)` with `console-event-handler` and calls `osi::ExitProcess` with exit code 80.

The scheduler maintains the *run queue*, a queue of ready-to-run processes, and the *sleep queue*, a queue of sleeping processes. Both are ordered by increasing precedence and preserve the order of insertion for processes with the same precedence. For the run queue, each process has precedence 0 in order to implement round-robin scheduling. For the sleep queue, each process uses its wake time as the precedence.

When the run queue is empty, the event-loop process calls `osi::GetCompletionPacket` with a non-zero timeout based on the first entry in the sleep queue to avoid busy waiting. When the event-loop process finishes processing all completion packets, it places itself at the end of the run queue.

Concern: Some process may starve another process. **Mitigation:** The run queue is managed with round-robin scheduling to prevent starvation. The event-loop process does not starve other processes because it drains the completion queue without causing new completion packets to be enqueued.

The finalizer process runs the finalizers registered via `add-finalizer`. These finalizers typically close operating system interface handles to objects that are no longer accessible. **Concern:** Ill-behaved finalizers may cause memory and handle leaks. **Mitigation:** Finalizers are designed to execute quickly without failing. Typical finalizers guard against errors when closing handles. If the finalizer process exits with reason *r*, the system logs the event `#(finalizer-process-terminated r)` with `console-event-handler` and calls `osi::ExitProcess` with exit code 80.

Once the finalizer process runs all the finalizers, it waits until another garbage collection has occurred before running again. The system hooks the `collect` procedure so that it sends a wake-up message to the finalizer process every time a garbage collection occurs. When the finalizer receives the wake-up message, it pumps all other wake-up messages from its inbox, since there may have been more than one garbage collection since it last ran.

Asynchronous I/O operations for COM ports, named pipes, external operating system processes, files, console input, and TCP connections are implemented with custom binary ports so that they have the same interface as the system I/O procedures. The system I/O procedures are not used because they perform synchronous I/O. The custom port buffer size is set to 1024³ with `custom-port-buffer-size`. The custom binary port read and write procedures call `osi::ReadPort` and `osi::WritePort` with callbacks that send a message to the calling process, which waits until it receives the message.

Concern: Using a port from more than one process at the same time may cause errors including buffer corruption. **Mitigation:** The code is inspected for concurrent use of ports. Port visibility is typically limited to a single process.

For two-way communication ports, we use two custom ports: one exclusively for input, and one exclusively for output. We do not use custom input/output ports for two reasons. First, textual input/output ports created with `transcoded-port` are not safe to use from two concurrent processes because one transcoding buffer is used for both reading and writing. Second, the input side of a port is commonly used only by a reader process, and the output side of a port is commonly used

³1024 was chosen because Chez Scheme uses 1024 for the buffer size of buffered transcoded ports.

only by a writer process. Keeping the input and output sides separate prevents concurrent use. The underlying handle is closed when the output port is closed.

Concern: Failing to close a handle from the operating system interface that is no longer used causes resource leaks. **Mitigation:** An *osi-port guardian* and associated finalizer are used to identify and close inaccessible osi-ports using `osi::ClosePort`. A *directory-watcher guardian* and associated finalizer are used to identify and close inaccessible directory watchers. A *listener guardian* and associated finalizer are used to identify and close inaccessible TCP listeners. In all cases, interrupts are disabled around code that wraps handles and registers objects with guardians in order to prevent the current process from being killed during this critical time.

3.4 Programming Interface

3.4.1 Process Creation

<code>(spawn thunk)</code>	procedure
returns: a process	

The `spawn` procedure creates and returns a new process that executes *thunk*, a procedure of no arguments. The new process starts with `name = #f`, `sic = 0` (interrupts enabled), `winders = ()`, an `exception-state` that terminates the process on an unhandled exception, an empty `inbox`, `precedence = 0`, `flags = 0` (the process is not sleeping and does not trap exits), `links = ()`, `monitors = ()`, and `src = #f`.

<code>(spawn&link thunk)</code>	procedure
returns: a process	

Like `spawn`, the `spawn&link` procedure creates and returns a new process that executes *thunk*. In addition, it links the new process to the calling process.

3.4.2 Process Registration

<code>(register name process)</code>	procedure
returns: #t	

The `register` procedure adds `name → process` to the registrar and sets `process.name = name`. When a registered process exits, its registration is removed. If *name* is not a symbol, exception `#(bad-arg register name)` is raised. If *process* is not a process, exception `#(bad-arg register process)` is raised. If *process* is dead, exception `#(process-dead process)` is raised. If *process* is already registered to name *n*, exception `#(process-already-registered n)` is raised. If *name* is already registered to process *p*, exception `#(name-already-registered p)` is raised.

<code>(unregister name)</code>	procedure
returns: #t	

The `unregister` procedure removes `name → process` from the registrar and sets `process.name = #f`. If *name* is not registered, exception `#(bad-arg unregister name)` is raised.

(whereis *name*) **procedure**

returns: a process | #f

The **whereis** procedure returns the process associated with *name* or #f if *name* is not registered. If *name* is not a symbol, exception **#(bad-arg whereis *name*)** is raised.

(get-registered) **procedure**

returns: a list of registered process names

The **get-registered** procedure returns a list of currently registered process names from the registrar.

3.4.3 Process Termination, Links, and Monitors

(catch *e1 e2 ...*) **syntax**

expands to:

```
(call/1cc
  (lambda (return)
    (with-exception-handler
      (lambda (reason) (return `(EXIT ,reason)))
      (lambda () e1 e2 ...))))
```

The **catch** macro evaluates expressions *e1 e2 ...* in a dynamic context that traps exceptions. If no exception is raised, the return value is the value of the last expression. If exception *reason* is raised, **#(EXIT *reason*)** is returned.

(kill *process reason*) **procedure**

returns: #t

The **kill** procedure is used to terminate a process.

1. If *process* is not a process, exception **#(bad-arg kill *process*)** is raised.
2. If *process* has already exited, nothing happens.
3. If *reason* is **kill**, *process* is terminated with reason **killed**, even if it traps exits.
4. If *process* traps exits, message **#(EXIT *self reason*)** is sent to *process*, where *self* is the calling process.
5. If *process* does not trap exits and *reason* is **normal**, nothing happens.
6. Otherwise, *process* is terminated with *reason*.

(link *process*) **procedure**

returns: #t

The **link** procedure creates a bi-directional link between the calling process (*self*) and *process*. No more than one link can exist between two processes, but it is not an error to call **link** more than once on the same two processes.

1. If *process* is not a process, exception `#{bad-arg link process}` is raised.
2. If *process* is *self*, nothing happens.
3. If *process* has not exited, then if the two processes are already linked, nothing happens; otherwise, *self* is added to *process.links*, and *process* is added to *self.links*.
4. Otherwise, *process* has exited with reason *r* = *process.exception-state*.
 - (a) If *self* traps exits, message `#{EXIT process r}` is sent to *self*.
 - (b) If *self* does not trap exits and *reason* is `normal`, nothing happens.
 - (c) Otherwise, *self* is terminated with reason *r*.

`(unlink process)` **procedure**
returns: `#t`

The `unlink` procedure removes the bi-directional link if present between the calling process (*self*) and *process* by removing *self* from *process.links* and *process* from *self.links*. If *process* is not a process, exception `#{bad-arg unlink process}` is raised.

`(monitor process)` **procedure**
returns: a monitor

The `monitor` procedure creates and returns a new monitor *m* with `origin` = the calling process (*self*) and `target` = *process*. Unlike `link`, `monitor` can create more than one connection between the same processes. It adds *m* to *self.monitors* and *process.monitors*. When *process* exits or has already exited with reason *r*, the message `#{DOWN m process r}` is sent to *self*. If *process* is not a process, exception `#{bad-arg monitor process}` is raised.

`(demonitor monitor)` **procedure**
returns: `#t`

The `demonitor` procedure removes a *monitor* created by the calling process (*self*) from *self.monitors* and *monitor.target.monitors* if present. If *monitor* is not a monitor with `origin` = *self*, exception `#{bad-arg demonitor monitor}` is raised.

`(demonitor&flush monitor)` **procedure**
returns: `#t`

The `demonitor&flush` procedure provides a convenient way to demonitor and flush any remaining DOWN message from the calling process's `inbox`. It performs the following operations:

```
(demonitor monitor)
(receive (until 0 #t)
  [#(DOWN ,@monitor ,_ ,_) #t])
```

`(monitor? x)` **procedure**
returns: a boolean

The `monitor?` procedure determines whether or not the datum *x* is a monitor.

pattern	matches
<i>symbol</i>	itself
<i>number</i>	itself
<i>boolean</i>	itself
<i>character</i>	itself
<i>string</i>	itself
<i>bytevector</i>	itself
()	itself
(<i>p</i> ₁ . <i>p</i> ₂)	a pair whose car matches <i>p</i> ₁ and cdr matches <i>p</i> ₂
#(<i>p</i> ₁ ... <i>p</i> _{<i>n</i>})	a vector of <i>n</i> elements whose elements match <i>p</i> ₁ ... <i>p</i> _{<i>n</i>}
,_	any datum
, <i>variable</i>	any datum and binds a fresh <i>variable</i> to it
,@ <i>variable</i>	any datum <code>equal?</code> to the bound <i>variable</i>
,(<i>variable</i> <= <i>pattern</i>)	any datum that matches <i>pattern</i> and binds a fresh <i>variable</i> to it
'(<i>type</i> {, <i>field</i> [<i>field</i> <i>pattern</i>]} ...)	an instance of the tuple <i>type</i> , each <i>field</i> of which is bound to fresh variable <i>field</i> or matches the corresponding <i>pattern</i>

Figure 3.1: Pattern Grammar

3.4.4 Messages and Pattern Matching

The pattern matching syntax of Figure 3.1 provides a concise and expressive way to match data structures and bind variables to parts. The `receive`, `match`, and `match-let*` macros use this pattern language. The implementation uses two structurally recursive passes over the pattern. The first pass checks the pattern syntax including tuple types and field names and accumulates the list of pattern variables. This list enables it to check for duplicate pattern variables. The second pass emits code that matches the input against the pattern left to right.

(send destination message) **procedure**
returns: *message*

The `send` procedure sends *message* to a process or registered name, *destination*. If *destination* is not a process or registered name, exception `#(bad-arg send destination)` is raised. If *destination* has exited, nothing else happens. Otherwise, *message* is added to the end of *destination.inbox*. If *destination* is sleeping, it is awakened. If *destination* is not on the run queue, it is placed on the run queue with precedence 0.

(receive **syntax**
 [(*after* *timeout* *t1* *t2* ...) | (*until* *time* *t1* *t2* ...)]
 (<pattern> [(*guard* *g*)] *b1* *b2* ...)
 ...)

returns: the value of the last evaluated expression

The `receive` macro examines each message *m* in the calling process's `inbox` by testing it against each pattern and optional guard. Each guard expression *g* is evaluated in the scope of its associated pattern variables. When *g* returns `#f`, *m* fails to match that clause. For the first pattern and guard

that matches *m*, *m* is removed from **inbox**, and the expressions *b1 b2 ...* are evaluated in the scope of its pattern variables. If *m* fails to match all patterns, the examination continues with the next message in **inbox**. When all messages have been examined, the calling process waits with its **src** field set to the source location of the **receive** macro if available. The process awakens when a new message or the time specified by the optional **after** or **until** clause arrives. If a new message arrives before the timeout, the examination process continues as before. Otherwise, the timeout expressions *t1 t2 ...* are evaluated.

The optional **after** clause specifies a *timeout* in milliseconds from the time at which control enters the **receive** macro. Similarly, the optional **until** clause specifies a clock *time* in milliseconds as measured by **erlang:now**. In addition, *timeout* and *time* can be **infinity** to indicate no timeout. If *t = timeout* or *time* is not a non-negative exact integer or **infinity**, exception **#(timeout-value t src)** is returned, where *src* is the source location of the **receive** macro if available.

See Figure 3.1 for the pattern grammar.

```
(match exp                                     syntax
  (<pattern> [(guard g)] b1 b2 ... )
  ...)
```

returns: the value of the last expression *b1 b2 ...* for the matched pattern

The **match** macro evaluates *exp* once and tests its value *v* against each pattern and optional guard. Each guard expression *g* is evaluated in the scope of its associated pattern variables. When *g* returns **#f**, *v* fails to match that clause. For the first pattern and guard that matches *v*, the expressions *b1 b2 ...* are evaluated in the scope of its pattern variables. If *v* fails to match all patterns, exception **#(bad-match v src)** is raised, where *src* is the source location of the **match** clause if available.

See Figure 3.1 for the pattern grammar.

```
(match-let* ([<pattern> [(guard g)] exp]       syntax
  ... )
  b1 b2 ...)
```

returns: the value of the last expression *b1 b2 ...*

The **match-let*** macro evaluates each *exp* in the order specified and matches its value against its pattern and guard. The pattern variables of each clause extend the scope of its guard expression *g* and all subsequent pattern clauses and body expressions *b1 b2 ...*. The **match-let*** macro returns the value of the last body expression. If any pattern fails to match or any *g* returns **#f**, exception **#(bad-match v src)** is raised, where *v* is the datum that failed to match the pattern or guard at source location *src* if available.

See Figure 3.1 for the pattern grammar.

3.4.5 Process Properties

```
self                                           syntax
```

returns: the current process

The **self** macro uses **identifier-syntax** to expand into code that retrieves the global self variable's top-level value. The global variable cannot be used directly because library bindings are immutable.

(process? <i>x</i>)	procedure
returns: a boolean	

The `process?` procedure determines whether or not the datum *x* is a process.

(process-id [<i>process</i>])	procedure
returns: the process id	

The `process-id` procedure returns *process.id*, where *process* defaults to `self`. If *process* is not a process, exception `#{bad-arg process-id process}` is raised.

process-trap-exit	parameter
value: boolean	

The `process-trap-exit` parameter specifies whether or not the calling process traps exit signals as messages. Processes start with this parameter set to `#f`.

(pps [<i>op</i>])	procedure
returns: unspecified	

The `pps` procedure prints information about all processes to *op*, which defaults to the current output port. If *op* is not an output port, exception `#{bad-arg pps op}` is raised.

3.4.6 Miscellaneous

(add-finalizer <i>finalizer</i>)	procedure
returns: unspecified	

The `add-finalizer` procedure adds *finalizer* to the global list of finalizers. *finalizer* is a procedure of no arguments that runs in the finalizer process after garbage collections. If *finalizer* is not a procedure, exception `#{bad-arg add-finalizer finalizer}` is raised.

(bad-arg <i>who arg</i>)	procedure
returns: never	

The `bad-arg` procedure raises exception `#{bad-arg who arg}`.

(console-event-handler <i>event</i>)	procedure
returns: unspecified	

The `console-event-handler` procedure prints an *event* to the console. It is used when the event manager is not available. It disables interrupts so that it can be called from multiple processes safely. The output is designed to be machine readable. The output looks like this:

```
Date: Fri Aug 06 11:54:59 2010
Timestamp: 1281110099144
Event: event
```

The date is the local time from the `date-and-time` procedure, the timestamp is the clock time from `erlang:now`, and *event* is printed as with `write`.

<code>(dbg)</code>	procedure
<code>(dbg <i>id</i>)</code>	
<code>(dbg <i>base proc</i>)</code>	
returns: see below	

The `dbg` procedure is used to debug processes that exit with a continuation condition.

`(dbg)` prints to the current output port the process id and exception message for each process that exited with a continuation condition.

`(dbg id)` enters the interactive debugger using the exception state of process *id*. If process *id* does not exist or did not exit with a continuation condition, the following message is printed: “Nothing to debug.”

`(dbg base proc)` folds over the processes that exited with a continuation condition and calls *proc* with the process id, process exception state, and the accumulator value (initially *base*). It returns the final accumulator value.

<code>(erlang:now)</code>	procedure
returns: the current clock time in milliseconds	

The `erlang:now` procedure calls `osi::GetTickCount` to return the number of milliseconds in UTC since the UNIX epoch January 1, 1970.

<code>(make-process-parameter <i>initial</i> [<i>filter</i>])</code>	procedure
returns: a process parameter procedure	

The `make-process-parameter` procedure creates a parameter procedure *p* that provides per-process, mutable storage using a weak eq-hashtable mapping processes to values. Calling *p* with no arguments returns the current value of the parameter for the calling process, and calling *p* with one argument sets the value of the parameter for the calling process. The *filter*, if present, is a procedure of one argument that is applied to the *initial* and all subsequent values. If *filter* is not a procedure, exception `#(bad-arg make-process-parameter filter)` is raised.

The following system parameters are not process safe and have been redefined to use `make-process-parameter`. `custom-port-buffer-size`, `exit-handler`, `pretty-initial-indent`, `pretty-line-length`, `pretty-maximum-lines`, `pretty-one-line-limit`, `pretty-standard-indent`, `print-brackets`, `print-char-name`, `print-gensym`, `print-graph`, `print-length`, `print-level`, `print-precision`, `print-radix`, `print-record`, `print-unicode`, `print-vector-length`, `reset-handler`, and `waiter-prompt-and-read`.

<code>(on-exit <i>finally</i> <i>b1</i> <i>b2</i> ...)</code>	syntax
expands to:	
<code>(dynamic-wind</code>	
<code>void</code>	
<code>(lambda () <i>b1</i> <i>b2</i> ...)</code>	
<code>(lambda () <i>finally</i>)</code>	

The `on-exit` macro executes the body expressions *b1 b2 ...* in a dynamic context that executes the *finally* expression whenever control leaves the body.

(profile-me) **procedure**

returns: unspecified

The **profile-me** procedure does nothing but provide a place-holder for the system profiler to count the call site. When profiling is turned off, **(profile-me)** expands to **(void)**, and the system optimizer eliminates it.

3.4.7 Tuples

For users of the concurrency model, a *tuple* is a container of named, immutable fields implemented as a vector whose first element is the tuple name and remaining elements are the fields. Each tuple definition is a macro that provides all tuple operations using field names only, not field indices. The macro makes it easy to copy a tuple without having to specify the fields that don't change. We decided not to use the Scheme record facility because it does not provide name-based constructors, copy operators, or convenient serialization.

(define-tuple name field ...) **syntax**

expands to: a macro definition of *name* described below

The **define-tuple** macro defines a macro for creating, copying, identifying, and accessing tuple type *name*. *name* and *field ...* must be identifiers. No two field names can be the same. The following field names are reserved: **make**, **copy**, **copy***, and **is?**.

(name make [field value] ...) **syntax**

returns: a new instance of tuple type *name* with *field = value ...*

The **make** form creates a new instance of the tuple type *name*. *field* bindings may appear in any order. All fields from the tuple definition must be specified.

(name field instance) **syntax**

returns: *instance.field*

The field accessor form retrieves the value of the specified *field* of *instance*. If *r = instance* is not a tuple of type *name*, exception **#(bad-tuple name r src)** is raised, where *src* is the source location of the field accessor form if available.

(name field) **syntax**

returns: a procedure that, given *instance*, returns *instance.field*

The **(name field)** form expands to **(lambda (instance) (name field instance))**.

(name open instance [prefix] (field ...)) **syntax**

expands to: definitions for *field ...* or *prefixfield ...* described below

The **open** form defines identifier syntax for each specified *field* so that a reference to *field* expands to **(name field r)** where *r* is the value of *instance*. If *r* is not a tuple of type *name*, exception **#(bad-tuple name r src)** is raised, where *src* is the source location of the **open** form if available. The **open** form is equivalent to the following, except that it checks the tuple type only once:

```
(begin
  (define instance instance)
  (define-syntax field (identifier-syntax (name field instance)))
  ...)
```

The `open` form introduces definitions only for fields listed explicitly in (*field* ...). If the optional *prefix* identifier is supplied, `open` produces a definition for *prefixfield* rather than *field* for each *field* specified.

```
(name copy instance [field value] ...) syntax
returns: a new instance of tuple type name with field = value ... and remaining fields copied from instance
```

The `copy` form creates a copy of *instance* except that each specified *field* is set to the associated *value*. If *r* = *instance* is not a tuple of type *name*, exception `#(bad-tuple name r src)` is raised, where *src* is the source location of the `copy` form if available. *field* bindings may appear in any order.

```
(name copy* instance [field value] ...) syntax
returns: a new instance of tuple type name with field = value ... and remaining fields copied from instance
```

The `copy*` form is like `copy` except that, within the *value* expressions, each specified *field* is bound to an identifier macro that returns the value of *instance.field*. If *r* = *instance* is not a tuple of type *name*, exception `#(bad-tuple name r src)` is raised, where *src* is the source location of the `copy*` form if available. The `copy*` form is equivalent to the following, except that it checks the tuple type only once:

```
(let ([instance instance])
  (name open instance (field ...))
  (name copy instance [field value] ...))
```

```
(name is? x) syntax
returns: a boolean
```

The `is?` form determines whether or not the datum *x* is an instance of tuple type *name*.

3.4.8 I/O

```
(make-utf8-transcoder) procedure
returns: a UTF-8 transcoder
```

The `make-utf8-transcoder` procedure creates a UTF-8 transcoder with end-of-line style `none` and error-handling mode `replace`.

```
(binary->utf8 bp) procedure
returns: a transcoded textual port wrapping bp
```

The `binary->utf8` procedure takes a binary port *bp* and returns a textual port wrapping *bp* using `transcoded-port` and `(make-utf8-transcoder)`. The original port *bp* is marked closed so that it cannot be used except through the associated textual port.

(force-close-output-port *op*) **procedure**
returns: unspecified

The **force-close-output-port** procedure is used to close a custom output port, even if it has unflushed output that would otherwise cause it to fail to close. If *op* is not already closed, **force-close-output-port** tries to close it with **(close-output-port *op*)**. If it fails with exception **#(io-error *name* WritePort *errno*)**, the output buffer is cleared with **(clear-output-port *op*)**, and **(close-output-port *op*)** is called again.

(io-error *name who errno*) **procedure**
returns: never

The **io-error** procedure raises exception **#(io-error *name who errno*)**. The string *name* identifies the file, pipe, COM port, process, console input, or TCP port. The symbol *who* specifies the procedure that raised an error, and the number *errno* specifies the error code. The **read-osi-port** procedure raises this exception with *who*=**ReadPort**, and the **write-osi-port** procedure raises it with *who*=**WritePort**.

(connect-usb *device-name read-address write-address*) **procedure**
returns: two values: a binary input port and a binary output port

The **connect-usb** procedure calls **osi::ConnectWinUSB** to establish a USB connection with the given *device-name*. It returns a custom binary input port that reads from the *read-address* endpoint and a custom binary output port that writes to the *write-address* endpoint. These ports do not track or report position. The operating system handle is closed when the output port is closed, not when the input port is closed, and the underlying *osi-port* is registered with the *osi-port* guardian.

(create-server-pipe *name callback*) **procedure**
returns: two values: a binary input port and a binary output port

The **create-server-pipe** procedure calls **osi::CreateServerPipe(*full-name*, *callback*)** to create a server pipe with *full-name*=**\\.\pipe***name*. It returns a custom binary input port that reads from the pipe and a custom binary output port that writes to the pipe. These ports do not track or report position. The operating system handle is closed when the output port is closed, not when the input port is closed, and the underlying *osi-port* is registered with the *osi-port* guardian.

The *callback* procedure is called with two arguments, *count* and *errno*, from the event loop after a client pipe has connected. Before this occurs, I/O operations on the server pipe will fail with error code 536.

If **osi::CreateServerPipe** returns error pair (*who* . *errno*), exception **#(io-error *full-name who errno*)** is raised.

(create-client-pipe *name*) **procedure**
returns: two values: a binary input port and a binary output port

The **create-client-pipe** procedure calls **osi::CreateClientPipe(*full-name*)** to create a client pipe connection to *full-name*=**\\.\pipe***name*. It returns a custom binary input port that reads from the pipe and a custom binary output port that writes to the pipe. These ports do not track or report position. The operating system handle is closed when the output port is closed, not when the input port is closed, and the underlying *osi-port* is registered with the *osi-port* guardian.

If `osi::CreateClientPipe` returns error pair (*who* . *errno*), exception `#(io-error full-name who errno)` is raised.

(create-watched-process *command-line callback*) **procedure**
returns: three values: a process handle, a binary input port, and a binary output port

The `create-watched-process` procedure calls `osi::CreateWatchedProcess(command-line, callback)` to spawn an operating system process. It returns a process handle that can be used with `osi::TerminateProcess`, a custom binary input port that reads from the standard output and standard error of the process, and a custom binary output port that writes to the standard input of the process. These ports do not track or report position. The operating system handle for the input port is closed when the input port is closed, and likewise for the output port. Both underlying `osi`-ports are registered with the `osi`-port guardian.

The *callback* procedure is called with two arguments, the process handle and the exit code, from the event loop after the operating system process exits.

If `osi::CreateWatchedProcess` returns error pair (*who* . *errno*), exception `#(create-watched-process-failed command-line who errno)` is raised.

(absolute-path *path root*) **procedure**
returns: the full, absolute path of *path* with respect to *root*

If *path* is an absolute path, the `absolute-path` procedure returns the full path of *path* using `osi::GetFullPath`. Otherwise, it returns the full path of `(path-combine root path)`.

(create-directory-path *path*) **procedure**
returns: *path*

The `create-directory-path` procedure creates directories as needed for the file *path*. It calls `osi::GetFullPath(path)` so that relative specifiers such as `..` work properly. It does not raise an exception if `osi::CreateDirectory` fails.

(create-file-port *name desired-access share-mode creation-disposition*) **procedure**
returns: an `osi`-port

The `create-file-port` procedure creates an `osi`-port by calling `osi::CreateFile(name, desired-access, share-mode, creation-disposition)`. The `osi`-port is registered with the `osi`-port guardian.

If `osi::CreateFile` returns error pair (*who* . *errno*), exception `#(io-error name who errno)` is raised.

The constants `GENERIC_READ` and `GENERIC_WRITE` can be used for *desired-access*.

The constants `FILE_SHARE_NONE`, `FILE_SHARE_READ`, `FILE_SHARE_WRITE`, and `FILE_SHARE_DELETE` can be used for *share-mode*.

The constants `CREATE_ALWAYS`, `CREATE_NEW`, `OPEN_ALWAYS`, `OPEN_EXISTING`, and `TRUNCATE_EXISTING` can be used for *creation-disposition*.

(read-osi-port *port bv start n fp*) **procedure**

returns: the number of bytes read

The `read-osi-port` procedure calls `osi::ReadPort` with the handle from the given osi-port *port*, bytevector buffer *bv*, starting 0-based buffer index *start*, maximum number of bytes to read *n*, starting 0-based file position *fp*, and a callback that fires in the event loop. The callback sends a message to the calling process, which waits until it receives it. If the read fails, exception `#(io-error name ReadPort errno)` is raised, where *name* is the name of *port* and *errno* is the error number. Otherwise, the number of bytes read is returned.

Error codes 38 (end of file), 109 (pipe closed), 995 (thread exit), and 10053 (connection aborted) are treated as end of file and cause a 0 to be returned.

(write-osi-port *port bv start n fp*) **procedure**

returns: the number of bytes written

The `write-osi-port` procedure calls `osi::WritePort` with the handle from the given osi-port *port*, bytevector buffer *bv*, starting 0-based buffer index *start*, maximum number of bytes to write *n*, starting 0-based file position *fp*, and a callback that fires in the event loop. The callback sends a message to the calling process, which waits until it receives it. If the write fails, exception `#(io-error name WritePort errno)` is raised, where *name* is the name of *port* and *errno* is the error number. Otherwise, the number of bytes written is returned.

(close-osi-port *port*) **procedure**

returns: unspecified

The `close-osi-port` procedure closes osi-port *port* using `osi::ClosePort`. If *port* has already been closed, `close-osi-port` does not raise an exception.

(get-file-size *port*) **procedure**

returns: the number of bytes in the file associated with osi-port *port*

The `get-file-size` procedure calls `osi::GetFileSize` to return the number of bytes in the file associated with osi-port *port*.

If `osi::GetFileSize` returns error pair (*who* . *errno*), exception `#(io-error filename who errno)` is raised.

(create-file *name desired-access share-mode creation-disposition type*) **procedure**

returns: a custom file port

The `create-file` procedure creates a custom file port by calling `(create-file-port name desired-access share-mode creation-disposition)`. The custom port supports both getting and setting the file position. The particular type of custom port returned is determined by *type*:

- **binary-input:** a binary input port
- **binary-output:** a binary output port
- **input:** a textual input port wrapping a binary input port with `binary->utf8`
- **output:** a textual output port wrapping a binary output port with `binary->utf8`

- **append**: a textual output port wrapping a binary output port with `binary->utf8` and initially positioned at the end of the file (determined by `osi::GetFileSize`)

If *type* is any other value, exception `#{bad-arg create-file type}` is raised.

(find-files *spec*) **procedure**

returns: a list of files and directories that match *spec*

The `find-files` procedure calls `osi::FindFiles` with a callback that fires in the event loop. The callback sends a message to the calling process, which waits until it receives the message and returns the list of files and directories that match *spec*. If `osi::FindFiles` or the message indicates error pair (*who* . *errno*), exception `#{find-files-failed spec who errno}` is raised.

(hook-console-input) **procedure**

returns: unspecified

The `hook-console-input` procedure replaces the system console input port, which uses synchronous I/O, with a custom textual input port that uses asynchronous I/O. It builds a custom binary input port with `osi::OpenConsole`, wraps it with `binary->utf8`, and sets the result as the `console-input-port`, `current-input-port`, and the system internal `$console-input-port`. It does nothing after it has been called once.

(move-file *old-pathname* *new-pathname* [*option*]) **procedure**

returns: `#t`

The `move-file` procedure calls `osi::MoveFile`. *option* may be either `error` or `replace`. A value of `replace` causes *new-pathname* to be replaced if it exists. The default is `error`.

If *old-pathname* is not a string, exception `#{bad-arg move-file old-pathname}` is raised. If *new-pathname* is not a string, exception `#{bad-arg move-file new-pathname}` is raised. If *option* is not `error` or `replace`, exception `#{bad-arg move-file option}` is raised.

(open-file-to-append *name*) **procedure**

returns: a custom file port

The `open-file-to-append` procedure calls `(create-file name GENERIC_WRITE FILE_SHARE_READ OPEN_ALWAYS 'append)`.

(open-file-to-read *name*) **procedure**

returns: a custom file port

The `open-file-to-read` procedure calls `(create-file name GENERIC_READ FILE_SHARE_READ OPEN_EXISTING 'input)`.

(open-file-to-replace *name*) **procedure**

returns: a custom file port

The `open-file-to-replace` procedure calls `(create-file name GENERIC_WRITE FILE_SHARE_READ CREATE_ALWAYS 'output)`.

(open-file-to-write *name*) **procedure**

returns: a custom file port

The `open-file-to-write` procedure calls `(create-file name GENERIC_WRITE FILE_SHARE_READ CREATE_NEW 'output)`.

(open-utf8-bytevector *bv*) **procedure**

returns: a transcoded textual input port wrapping *bv*

The `open-utf8-bytevector` procedure calls `(binary->utf8 (open-bytevector-input-port bv))`.

(path-combine *path*₁ *path*₂ ...) **procedure**

returns: the string combining the paths

The `path-combine` procedure appends one or more paths, inserting the `\` character between each pair of paths when the first path of the pair does not end with it.

(get-source-offset *ip*) **procedure**

returns: an exact nonnegative integer

The `get-source-offset` procedure takes a binary input port *ip* that supports `port-position`, skips over the `#!interpreter-directive` line, if any, and returns the resulting `port-position`.

(get-datum/annotations-all *ip sfd bfp*) **procedure**

returns: a list of annotated objects

The `get-datum/annotations-all` procedure takes a textual input port *ip*, a source-file descriptor *sfd*, and an exact nonnegative integer *bfp* representing the character position of the next character to be read from *ip*. The procedure returns a list of the annotated objects, in order, obtained by repeatedly calling `get-datum/annotations` with the advancing *bfp*, until *ip* reaches the end of file.

(with-sfd-source-offset *name handler*) **procedure**

returns: see below

The `with-sfd-source-offset` procedure takes a filename *name* and returns the result of calling the procedure *handler* with three arguments: *ip*, a textual port transcoded with the UTF-8 transcoder described in `binary->utf8`, *sfd*, a source-file descriptor that refers to *name*, and *source-offset*, the value returned by `get-source-offset`. Before returning, `with-sfd-source-offset` closes the textual port.

(read-bytevector *name contents*) **procedure**

returns: a list of annotations

The `read-bytevector` procedure takes a filename *name* and *contents* bytevector and returns a list of annotations read using `get-datum/annotations` from the *contents* bytevector transcoded with the UTF-8 transcoder described in `binary->utf8`.

(read-file *name*) **procedure**

returns: a bytevector with the contents of *name*

The `read-file` procedure calls `(create-file-port name GENERIC_READ FILE_SHARE_READ OPEN_EXISTING)` to open the file *name* and returns the contents as a bytevector.

(watch-directory path subtree? callback) **procedure**

returns: a directory watcher

The `watch-directory` procedure calls `osi::WatchDirectory` to track changes in the directory *path* and returns a directory watcher that is registered with the directory-watcher guardian. When *subtree?* is true, the directory tree rooted at *path* is watched. The change filter is configured to detect creating, renaming, and deleting files and directories as well as changes to the size and last-write time of files. Changes to file attributes, last-access and creation times, and security attributes are not included.

Every time changes are detected, `(callback ((action . filename) ...))` is executed in the event loop. The *action* number is defined as follows:

<i>action</i>	Meaning
1	<i>filename</i> was added to the directory.
2	<i>filename</i> was removed from the directory.
3	<i>filename</i> was modified.
4	<i>filename</i> was renamed, and <i>filename</i> is the old name.
5	<i>filename</i> was renamed, and <i>filename</i> is the new name.

When the directory watcher stops, it executes `(callback errno)` in the event loop, and no more callbacks from this watcher will be executed. An *errno* of zero indicates a normal close from `close-directory-watcher`. A non-zero *errno* indicates an error from `ReadDirectoryChangesW`.

If `osi::WatchDirectory` returns error pair `(who . errno)`, exception `#(watch-directory-failed path who errno)` is raised.

(close-directory-watcher watcher) **procedure**

returns: unspecified

The `close-directory-watcher` procedure uses `osi::CloseDirectoryWatcher` to close the given directory *watcher*. If the watcher has not already stopped, the associated callback will be executed at least one more time. If *watcher* is not a directory watcher, exception `#(bad-arg close-directory-watcher watcher)` is raised. If *watcher* has already been closed, `close-directory-watcher` does not raise an exception.

(directory-watcher-path watcher) **procedure**

returns: the *path* field of *watcher*

The `directory-watcher-path` procedure returns the *path* of the given directory *watcher*.

(listen-tcp address port-number process) **procedure**

returns: a TCP listener

The `listen-tcp` procedure calls `osi_listen_tcp` to create a TCP listener on the given *address* and *port-number* and returns a TCP listener that is registered with the listener guardian.

For each accepted connection, the message `#(accept-tcp listener ip op)` is sent to *process*, where *ip* is the input port and *op* is the output port.

For each failed connection, the message `#(accept-tcp-failed listener who errno)` is sent to *process*, where *who* and *errno* specify the error.

The *address* is a dotted quad IPv4 address or an IPv6 address. Use `":: "` to listen on all IPv4 and IPv6 interfaces. Use `"0.0.0.0"` to listen on all IPv4 interfaces. Otherwise, it listens on the given *address* only. If *address* is not a string, exception `#(bad-arg listen-tcp address)` is raised.

If *port-number* is zero, the operating system will choose an available port number, which can be queried with `listener-port-number`. If *port-number* is not a fixnum between 0 and 65535 inclusive, exception `#(bad-arg listen-tcp port-number)` is raised.

If `osi_listen_tcp` returns error pair `(who . errno)`, exception `#(listen-tcp-failed address port-number who errno)` is raised.

<code>(close-tcp-listener listener)</code>	procedure
returns: unspecified	

The `close-tcp-listener` procedure closes a TCP *listener* using `osi::CloseTCPListener`. If *listener* is not a TCP listener, exception `#(bad-arg close-tcp-listener listener)` is raised. If *listener* has already been closed, `close-tcp-listener` does not raise an exception.

<code>(listener-address listener)</code>	procedure
returns: the <i>address</i> field of <i>listener</i>	

The `listener-address` procedure returns the *address* of the given TCP *listener*.

<code>(listener-port-number listener)</code>	procedure
returns: the <i>port-number</i> field of <i>listener</i>	

The `listener-port-number` procedure returns the *port-number* of the given TCP *listener*.

<code>(connect-tcp hostname port-spec [process])</code>	procedure
returns: see below	

The `connect-tcp` procedure calls `osi::ConnectTCP` with a callback that fires in the event loop when the TCP connection to *hostname* on *port-spec* is established or fails to be established. The *port-spec* may be a fixnum port number or a string service name such as `"http"` . The callback creates a custom binary input port that reads from the new connection and a custom binary output port that writes to the new connection. These ports do not track or report position. The operating system handle is closed when the output port is closed, not when the input port is closed, and the underlying `osi-port` is registered with the `osi-port` guardian. The callback sends one of the following messages to *process*, which defaults to `self`:

- `#(connect-tcp hostname port-spec ip op)`, where *ip* is the input port and *op* is the output port

- `#{connect-tcp-failed hostname port-spec who errno}`, where *who* and *errno* specify the error from the callback to `osi::ConnectTCP`

When *process* is present, `connect-tcp` returns `#t` without waiting for the callback.

When *process* is absent, `connect-tcp` waits to receive a message from the callback. If it receives the `connect-tcp` message, it returns two values, *ip* and *op*. If it receives the `connect-tcp-failed` message, it raises exception `#{connect-tcp-failed hostname port-spec who errno}`.

If *hostname* is not a string, exception `#{bad-arg connect-tcp hostname}` is raised. If *port-spec* is not a fixnum between 0 and 65535 inclusive or a string, exception `#{bad-arg connect-tcp port-spec}` is raised.

3.4.9 Queues

A queue is represented as a pair of lists, (*in* . *out*). The *out* list contains the first elements of the queue, and the *in* list contains the last elements of the queue in reverse. This representation allows for O(1) amortized insertion and removal times. The implementation is based on the Erlang queue module [11].

<code>(queue:add x q)</code>	procedure
returns: a queue that adds <i>x</i> to the rear of <i>q</i>	

<code>(queue:add-front x q)</code>	procedure
returns: a queue that adds <i>x</i> to the front of <i>q</i>	

<code>(queue:drop q)</code>	procedure
returns: a queue without the first element of <i>q</i>	

<code>queue:empty</code>	syntax
returns: the empty queue	

<code>(queue:empty? q)</code>	procedure
returns: <code>#t</code> if <i>q</i> is a queue, <code>#f</code> otherwise	

<code>(queue:get q)</code>	procedure
returns: the first element of <i>q</i>	

3.4.10 Hash Tables

The implementation of functional hash tables is based on the Erlang dict module [9, 14].

<code>(ht:delete ht key)</code>	procedure
returns: a hash table formed by dropping any association of <i>key</i> from <i>ht</i>	

(ht:fold <i>ht f init</i>)	procedure
returns: see below	

The **ht:fold** procedure accumulates a value by applying f to each key/value association in ht and the accumulator, which is initially $init$. It can be defined recursively as follows, where n is the size of ht , and the result of **ht:fold** is F_n :

$$F_0 = init$$

$$F_i = (f \ key_i \ val_i \ F_{i-1}) \text{ for } 1 \leq i \leq n$$

(ht:is? <i>x</i>)	procedure
returns: #t if x is a hash table, #f otherwise	

(ht:keys <i>ht</i>)	procedure
returns: a list of the keys of ht	

(ht:make <i>hash-key equal-key? valid-key?</i>)	procedure
returns: an empty hash table	

The **ht:make** procedure returns an empty hash table.

The *hash-key* procedure takes a key and returns an exact integer. It must return the same integer for equivalent keys.

The *equal-key?* procedure takes two keys and returns a true value if they are equivalent and **#f** otherwise.

The *valid-key?* procedure takes a datum and returns a true value if it a valid key and **#f** otherwise.

(ht:ref <i>ht key default</i>)	procedure
returns: the value associated with key in ht , $default$ if none	

(ht:set <i>ht key val</i>)	procedure
returns: a hash table formed by associating key with val in ht	

(ht:size <i>ht</i>)	procedure
returns: the number of entries in ht	

3.4.11 Error Strings

current-exit-reason->english	parameter
value: a procedure of one argument that returns an English string	

The **current-exit-reason->english** parameter specifies the conversion procedure used by **exit-reason->english**. It defaults to **swish-exit-reason->english**.

(exit-reason->english <i>x</i>)	procedure
---	------------------

returns: a string in U.S. English

The `exit-reason->english` procedure converts an exit reason into an English string using the procedure stored in parameter `current-exit-reason->english`.

(swish-exit-reason->english *x*) **procedure**

returns: a string in U.S. English

The `swish-exit-reason->english` procedure converts an exit reason from Swish into an English string.

3.4.12 String Utilities

The string utilities below are found in the `(swish string-utils)` library. For regular expression support, see the `(swish pregexp)` library described in [22].

(ends-with? *s p*) **procedure**

returns: a boolean

The `ends-with?` procedure determines whether or not the string *s* ends with string *p* using case-sensitive comparisons.

(ends-with-ci? *s p*) **procedure**

returns: a boolean

The `ends-with-ci?` procedure determines whether or not the string *s* ends with string *p* using case-insensitive comparisons.

(format-rfc2822 *d*) **procedure**

returns: a string like “Thu, 28 Jul 2016 17:20:11 -0400”

The `format-rfc2822` procedure returns a string representation of the date object *d* in the form specified in Section 3.3 of RFC 2822 [21].

(join *ls separator* [*last-separator*]) **procedure**

returns: a string

The `join` procedure returns the string formed by displaying each of the elements of list *ls* separated by displaying *separator*. When *last-separator* is specified, it is used as the last separator.

(split *str separator*) **procedure**

returns: a list of strings

The `split` procedure divides the *str* string by the *separator* character into a list of strings, none of which contain *separator*.

(split-n *str separator n*) **procedure**

returns: a list of no more than n strings

The `split-n` procedure divides the *str* string by the *separator* character into a list of at most n strings. The last string may contain *separator*.

`(starts-with? s p)` **procedure**

returns: a boolean

The `starts-with?` procedure determines whether or not the string *s* starts with string *p* using case-sensitive comparisons.

`(starts-with-ci? s p)` **procedure**

returns: a boolean

The `starts-with-ci?` procedure determines whether or not the string *s* starts with string *p* using case-insensitive comparisons.

`(symbol-append . ls)` **procedure**

returns: a symbol

The `symbol-append` procedure returns the symbol formed by appending the symbols passed as arguments.

Chapter 4

Generic Server

The generic server provides an “empty” server, that is, a framework from which instances of servers can be built.
—Joe Armstrong [1]

4.1 Introduction

In a concurrent system, many processes need to access a shared resource or sequentially manipulate the state of the system. This is generally modeled using a client/server design pattern. To help developers build robust servers, a generic server (**gen-server**) implementation inspired by Erlang’s Open Telecom Platform is provided.

The principles of the generic server can be found in Joe Armstrong’s thesis [1] or *Programming Erlang—Software for a Concurrent World* [2]. Documentation for Erlang’s **gen_server** is available online [10]. Source code for the Erlang Open Telecom Platform can be found online [8]. The source code for **gen_server** is part of `stdlib` and can be found in `/lib/stdlib/src/gen_server.erl`.

4.2 Theory of Operation

A **gen-server** provides a consistent mechanism for programmers to create a process which manages state, timeout conditions, and failure conditions using functional programming techniques. A programmer uses **gen-server:start&link** and implements the callback API to instantiate particular behavior.

A generic server starts a new process, registers it as a named process, and invokes the `init` callback procedure while blocking the calling process.

Clients can then send messages to a server using the synchronous **gen-server:call**, the asynchronous **gen-server:cast**, or the raw `send` procedure. The **gen-server** framework will automatically process messages and dispatch them to `handle-call`, `handle-cast`, and `handle-info` respectively.

The **gen-server** framework code automatically interprets a `stop` return value from the callback API or an `EXIT` message from the process which created it as a termination request and calls `terminate`.

Unless the termination reason is either `normal` or `shutdown`, generic servers use `event-mgr:notify` to report the termination.

Erlang's `gen_server` supports timeouts during `gen_server:start` and `gen_server:start&link`. In order to simplify the startup code, we have not implemented this feature. Timeouts while running the `init` callback may cause resources to be stranded until the garbage collector can clean them up. Timeouts during initialization should be considered carefully.

4.3 Programming Interface

(gen-server:start&link name arg ...) **syntax**

returns: `#(ok pid) | #(error reason) | ignore`

name: a symbol for a registered server or `#f` for an anonymous server

arg: any Scheme datum

`gen-server:start&link` spawns the server process, links to the calling process, registers the server process as *name*, and calls `(init arg ...)` within that process. To ensure a synchronized startup procedure, `gen-server:start&link` does not return until `init` has returned.

This macro uses the current scope to capture the callback functions `init`, `handle-call`, `handle-cast`, `handle-info`, and `terminate`.

Attempting to register a name that already exists results in `#(error #(name-already-registered pid))`, where *pid* is the existing process.

The return value of `gen-server:start&link` is propagated from the `init` callback.

An `init` which returns `#(ok state [timeout])` will yield `#(ok pid)` where *pid* is the newly created process.

An `init` which returns `#(stop reason)` or exits with *reason* will terminate the process and yield `#(error reason)`.

An `init` which returns `ignore` will terminate the process and yield `ignore`. This value is useful to inform a supervisor that the `init` procedure has determined that this server is not necessary for the system to operate.

An `init` which returns *other* values will terminate the process and yield `#(error #(bad-return-value other))`.

(gen-server:start name arg ...) **syntax**

returns: `#(ok pid) | #(error error) | ignore`

`gen-server:start` behaves the same as `start&link` except that it does not link to the calling process.

(gen-server:call server request [timeout]) **procedure**

returns: *reply*

server: process or registered name

request: any Scheme datum

timeout: non-negative exact integer in milliseconds or `infinity`, defaults to 5000

`gen-server:call` sends a synchronous *request* to *server* and waits for a *reply*. The server processes the request using `handle-call`.

Failure to receive a reply causes the calling process to exit with reason `#{timeout #(gen-server call (server request))}` if no timeout is specified, or `#{timeout #(gen-server call (server request timeout))}` if a timeout is specified. If the caller catches the failure and continues running, the caller must be prepared for a possible late reply from the server.

`gen-server:call` exits if the server terminates while the client is waiting for a reply. When that happens, the client exits for the same reason as the server.

<code>(gen-server:cast server request)</code>	procedure
returns: <code>ok</code>	

server: process or registered name
request: any Scheme datum

`gen-server:cast` sends an asynchronous *request* to a *server* and returns `ok` immediately. When using `gen-server:cast` a client does not expect failures in the server to cause failures in the client; therefore, this procedure ignores all failures. The server will process the request using `handle-cast`.

<code>(gen-server:reply client reply)</code>	procedure
returns: <code>ok</code>	

client: a *from* argument provided to the `handle-call` callback
reply: any Scheme datum

`gen-server:reply` can be used by a server to explicitly send a *reply* to a *client* that called `gen-server:call`.

In some situations, a server cannot reply immediately to a client. In such cases, the *from* argument inside `handle-call` is stored, and `no-reply` is returned. Later, `gen-server:reply` can be called using that *from* value as *client*. The *reply* is the return value of the `gen-server:call` in this case.

<code>(gen-server:debug server server-options client-options)</code>	procedure
returns: <code>ok</code>	

server: process or registered name
server-options: (`[message]` `[state]` `[reply]`) | `#f`
client-options: (`[message]` `[reply]`) | `#f`

`gen-server:debug` sets the debugging mode of *server*. The *server-options* argument specifies the logging of calls in the server. When *server-options* is `#f`, server logging is turned off. Otherwise, server logging is turned on, and *server-options* is a list of symbols specifying the level of detail. In logging mode, the *server* sends a `<gen-server-debug>` event for each call to `handle-call`, `handle-cast`, and `handle-info`. The *message* field is populated when `message` is in *server-options*, the *state* field is populated when `state` is in *server-options*, and the *reply* field is populated when `reply` is in *server-options*.

Similarly, the *client-options* argument specifies the logging of client calls to *server* with `gen-server:call`. When *client-options* is `#f`, client logging is turned off. Otherwise, client logging is turned on, and *client-options* is a list of symbols specifying the level of detail. In logging mode, `gen-server:call` sends a `<gen-server-debug>` event. The *message* field is populated when `message` is in *client-options*, and the *reply* field is populated when `reply` is in *client-options*.

<code>(define-state-tuple <i>name field ...</i>)</code>	syntax
---	---------------

This form defines a tuple type using `(define-tuple name field ...)` and defines a new syntactic form `$state`. `$state` provides a succinct syntax for the `state` variable.

`$state` transforms `($state op arg ...)` to `(name op state arg ...)` where `state` is a variable in the same scope as `$state`.

Given this definition:

```
(define-state-tuple <my-state> x y z)
```

The following code is equivalent:

```
(<my-state> copy state [x 2])  ($state copy [x 2])
(<my-state> x state)           ($state x)
(<my-state> y state)           ($state y)
(<my-state> z state)           ($state z)
```

There is no equivalent for constructing a state tuple because constructing a tuple does not require the `state` variable. The `(<my-state> make ...)` syntax must be used.

4.4 Published Events

All generic servers send the event manager the following event:

<gen-server-terminating>	event
---------------------------------------	--------------

timestamp: the time the event occurred
name: the name of the server
last-message: the last message received by the server
state: the last state passed into `terminate`
reason: the reason for termination

This event is fired after a successful call to `terminate` unless the reason for termination is `normal` or `shutdown`. If the `terminate` procedure exits with a new *reason*, the event contains the new *reason*.

<gen-server-debug>	event
---------------------------------	--------------

timestamp: the time the operation started
duration: the duration of the operation in milliseconds
type: 1 for `handle-call`, 2 for `handle-cast`, 3 for `handle-info`, 4 for `terminate`, 5 for a successful `gen-server:call`, and 6 for a failed `gen-server:call`
client: the client process or `#f`
server: the server process
message: the message sent to the server or `#f`
state: the state of the server when it received the message or `#f`
reply: the server's reply or `#f`

4.5 Callback Interface

A programmer implements the callback interface to define a particular server's behavior. All callback functions are called from within the server process.

The callback functions for gen-server processes are supposed to be *well-behaved functions*, i.e., functions that work correctly. The generation of an exception in a well-behaved function is interpreted as a failure [1].

When a callback function exits with a reason, `terminate` is called and the server exits.

When a callback function returns an unexpected *value*, `terminate` is called with the reason `#{bad-return-value value}`, and the server exits.

A callback may specify a *timeout* as a relative time in milliseconds up to one day, an absolute time in milliseconds (e.g., from `erlang:now`), or `infinity`. The default *timeout* is `infinity`. If the time period expires before another message is received, then a `timeout` message will be processed by `handle-info`.

Messages sent using `send`, including `#{EXIT pid reason}` and `#{DOWN monitor pid reason}`, are processed by `handle-info`.

The generic server framework will automatically interpret an EXIT message from the process which spawned it as a reason for termination. `terminate` will be called directly. `handle-info` will not be called. The server must use `(process-trap-exit #t)` to receive EXIT messages.

<code>(init arg ...)</code>	procedure
returns: <code>#{ok state [timeout]} #{stop reason} ignore</code>	

arg ...: the *arg ...* provided to `gen-server:start&link` or `gen-server:start`

state: any Scheme datum

timeout: relative time in milliseconds up to one day, absolute time in milliseconds
(e.g., from `erlang:now`), or `infinity` (default)

reason: any Scheme datum

`(init arg ...)` is called from a new server process and must complete before `gen-server:start&link` or `gen-server:start` returns.

A successful `init` returns `#{ok state [timeout]}`. The *state* is then maintained functionally by the generic server framework.

`init` may specify that server initialization failed by returning `#{stop reason}`. The server will then fail to start using this *reason*. `terminate` will not be called as the server has not properly started.

`init` may specify `ignore`. The server will then exit with reason `normal`, and `gen-server:start&link` will return `ignore`. This is used to inform a supervisor that the server is not necessary for the system to operate. `terminate` will not be called.

<code>(handle-call request from state)</code>	procedure
returns: <code>#{reply reply state [timeout]} #{no-reply state [timeout]} #{stop reason [reply] state}</code>	

request: the *request* provided to `gen-server:call`
from: `#{(client-process tag)}`
state: server state
reply: any Scheme datum
timeout: relative time in milliseconds up to one day, absolute time in milliseconds
(e.g., from `erlang:now`), or `infinity` (default)
reason: any Scheme datum

`handle-call` is responsible for processing a client *request* generated by `gen-server:call`.

`handle-call` may return `#{reply reply state [timeout]}` to indicate that *reply* is to be returned from `gen-server:call` to the caller. The server state will become *state*.

`handle-call` may return `#{no-reply state [timeout]}` to continue operation and to indicate that the caller of `gen-server:call` will continue to wait for a reply. The server state will become *state*. The server will need to use `gen-server:reply` and *from* to reply to the client.

`handle-call` may return `#{stop reason [reply] state}` to set a new *state*, then terminate the server with the given *reason*. If the optional *reply* is specified, it will be the return value of `gen-server:call`; otherwise, `gen-server:call` will exit with *reason*.

reply is any Scheme datum.

state is any Scheme datum.

reason is any Scheme datum.

(handle-cast request state)	procedure
returns: <code>#{(no-reply state [timeout])} #{(stop reason state)}</code>	

request: the *request* provided to `gen-server:cast`
state: server state
timeout: relative time in milliseconds up to one day, absolute time in milliseconds
(e.g., from `erlang:now`), or `infinity` (default)
reason: any Scheme datum

`handle-cast` is responsible for processing a client *request* generated by `gen-server:cast`.

`handle-cast` may return `#{(no-reply state [timeout])}` to continue operation. The server state will become *state*.

`handle-cast` may return `#{(stop reason state)}` to terminate the server with the given *reason*. The server state will become *state*.

(handle-info msg state)	procedure
returns: <code>#{(no-reply state [timeout])} #{(stop reason state)}</code>	

msg: `timeout` or a Scheme datum sent via `send`
state: server state
timeout: relative time in milliseconds up to one day, absolute time in milliseconds
(e.g., from `erlang:now`), or `infinity` (default)
reason: any Scheme datum

`handle-info` is responsible for processing timeouts and miscellaneous messages sent to the server via `send`.

`handle-info` may return `#{no-reply state [timeout]}` to continue operation. The server state will become *state*.

`handle-info` may return `#{stop reason state}` to terminate the server with the given *reason*. The server state will become *state*.

<code>(terminate reason state)</code>	procedure
returns: ignored	

reason: shutdown reason

state: server state

`terminate` is called when the server is about to terminate. It is responsible for cleaning up any resources that the server allocated. When it returns, the server exits for the given *reason*.

reason can be any reason specified by a stop return value `#{stop ...}`. When a supervision tree is terminating, *reason* will be `shutdown`.

The return value of `terminate` is ignored. Unless the termination reason is either `normal` or `shutdown`, the generic server framework uses `event-mgr:notify` to report the termination. The server then terminates for that reason.

If `terminate` exits with *reason*, then that reason is logged, and the server terminates with *reason*.

Chapter 5

Event Manager

5.1 Introduction

The event manager (`event-mgr`) is a gen-server that provides a single dispatcher for events within the system. It buffers events and dispatches them to the log handler and a collection of other event handlers. If the log handler fails, the event manager logs events directly to the console.

5.2 Theory of Operation

The event manager is a singleton process through which all events in the system are routed. Any component may notify the event manager that something has occurred by using `event-mgr:notify`. This model is illustrated in Figure 5.1.

The event manager is a registered process named `event-mgr`.

The event manager is created as part of the application's supervision hierarchy. It buffers incoming events during startup until `event-mgr:flush-buffer` is called. The buffered events are then sent to the current event handlers and the log handler. This provides the ability to log the startup details of processes, including the event manager itself.

The event handlers should not perform blocking operations, because they block the entire event manager.

If the log handler or its associated process fails, the event manager logs events to the console. If another event handler fails with some reason, the associated process is killed with the same reason. When the process associated with a handler terminates, the event manager removes it from the list.

state

<event-mgr-state>

tuple




Figure 5.1: Event flow

event-buffer: list of events to be processed (most recent first), or **#f** when buffering is disabled

log-handler: **<handler>** tuple or **#f**

handlers: list of **<handler>** tuples

<handler>	tuple
------------------------	--------------

proc: procedure of one argument, the event

owner: process that owns the handler

init The **init** procedure initializes the state of the gen-server. Event buffering is enabled.

The gen-server traps exits so that it can detect failure of event handler owner processes, as well as the **EXIT** message from the parent process.

terminate The **terminate** procedure flushes any pending events to the console using **do-notify**.

handle-call The **handle-call** procedure processes the following messages:

- **#(add-handler proc owner)**: Link to the *owner* process, add a handler to the state and return **ok**.

An invalid argument results in the following error reasons:

- **#(invalid-procedure proc)**

- `#{invalid-owner owner}`
- `flush-buffer`: Process the events in the buffer using `do-notify`, turn off buffering, and return `ok`.
- `#{set-log-handler proc owner}`: Link to the *owner* process, set the log handler of the state, and return `ok`.

An invalid argument results in the following error reasons:

- `log-handler-already-set`
- `#{invalid-procedure proc}`
- `#{invalid-owner owner}`

`handle-cast` The `handle-cast` procedure does not process any messages.

`handle-info` The `handle-info` procedure handles the following messages:

- `#{notify event}`: Process *event* using `do-notify`.
event is any Scheme datum.
- `#{EXIT pid _}`: Removes the log or other event handler associated with *pid*.

Internally, the `(do-notify event state)` procedure handles the processing of each *event* with respect to the current *state*. It evaluates the *state* in the following way:

- If the state is not buffering:
 1. Call each handler's *proc* with *event*. If it exits for some reason, kill the handler's *owner* with the same reason.
 2. If there is a log handler, call its *proc* with *event*. If it exits for some reason, unlink its *owner*, kill it with the same reason, log *event* to the console using `console-event-handler`, and remove the log handler from the state.
- Otherwise, buffer the event.

5.3 Programming Interface

`(event-mgr:start&link)` procedure
returns: `#{ok pid} | #{error reason}`

The `event-mgr:start&link` procedure creates a new `event-mgr` gen-server using `gen-server:start&link`.

The event manager is registered as `event-mgr`.

`(event-mgr:add-handler proc [owner])` procedure

returns: `ok` | `#(error reason)`

The `event-mgr:add-handler` procedure calls `(gen-server:call event-mgr #(add-handler proc owner))`.

proc is a procedure of one argument, the event. Failure in *proc* results in the event manager killing the *owner* process with the same failure reason. The handler is removed when the event manager receives an `EXIT` message from *owner*.

owner is a process. The default is the calling process.

(event-mgr:flush-buffer)	procedure
returns: <code>ok</code>	

The `event-mgr:flush-buffer` procedure calls `(gen-server:call event-mgr flush-buffer)`.

(event-mgr:notify event)	procedure
returns: <code>ok</code>	

The `event-mgr:notify` procedure sends message `#(notify event)` to registered process `event-mgr` if it exists. If `event-mgr` does not exist, it prints *event* using `console-event-handler`.

event is any Scheme datum.

Because the `gen-server` library uses `event-mgr:notify`, it is implemented there.

(event-mgr:set-log-handler proc owner)	procedure
returns: <code>ok</code> <code> #(error reason)</code>	

The `event-mgr:set-log-handler` procedure calls `(gen-server:call event-mgr #(set-log-handler proc owner))`.

proc is a procedure of one argument, the event. Failure in *proc* results in the event manager killing the *owner* process with the same failure reason. The log handler is removed when *proc* fails or the event manager receives an `EXIT` message from *owner*.

owner is a process.

Chapter 6

Gatekeeper

6.1 Introduction

The gatekeeper is a single gen-server named `gatekeeper` that manages shared resources using mutexes. Before a process uses a shared resource, it asks the gatekeeper to enter the corresponding mutex. When the process no longer needs the resource or terminates, it tells the gatekeeper to leave the mutex. A process may enter the same mutex multiple times, and it needs to leave the mutex the same number of times. The gatekeeper breaks deadlocks by raising an exception in one of the processes waiting for a mutex involved in a cyclic dependency chain.

The gatekeeper hooks system primitives `$cp0`, `$np-compile`, `pretty-print`, and `sc-expand` because they are not safe to be called from two processes at the same time (see the discussion of the global winders list in Section 3.3). The `$cp0` procedure uses resource `$cp0`, the `$np-compile` procedure uses resource `$np-compile`, and so forth.

6.2 Theory of Operation

state The gatekeeper state is a list of `<mutex>` tuples, each of which has the following fields:

- *resource*: resource compared for equality using `eq?`
- *process*: process that owns *resource*
- *monitor*: monitor of *process*
- *count*: number of times *process* has entered this mutex
- *waiters*: ordered list of *from* arguments from `handle-call` for processes that are waiting to enter this mutex

init The gatekeeper `init` procedure hooks the system primitives listed in the introduction so that they use `with-gatekeeper-mutex` with a timeout of one minute, and it sets the `current-expand` parameter to the hooked `sc-expand` procedure. The process traps exits so that `terminate` can unhook the system primitives when the process is shut down. It returns an empty list of `<mutex>` tuples.

terminate The gatekeeper **terminate** procedure unhooks the system primitives listed in the introduction and sets the **current-expand** parameter to the unhooked **sc-expand** procedure.

handle-call The gatekeeper **handle-call** procedure handles the following messages:

- **#(enter resource)**: Find $mutex \in state$ where $mutex.resource = resource$.
If no such $mutex$ exists, no-reply with **(enter-mutex resource from '() state)**.
If $mutex.process = from.process$, increment $mutex.count$, and reply **ok** with the updated state.
If **(deadlock? from.process mutex state)**, reply **#(deadlock resource)** with $state$.
Otherwise, add $from$ to the end of $mutex.waiters$, and no-reply with the updated state.
- **#(leave resource)**: Find $mutex \in state$ where $mutex.resource = resource$ and $mutex.process = from.process$.
If no such $mutex$ exists, reply **#(unowned-resource resource)** with $state$.
If $mutex.count > 1$, decrement $mutex.count$, and reply **ok** with the updated state.
Otherwise, reply **ok** with **(leave-mutex mutex state)**.

handle-cast The gatekeeper **handle-cast** procedure raises an exception on all messages.

handle-info The gatekeeper **handle-info** procedure handles the following message:

- **#(DOWN monitor _ _)**: Find $mutex \in state$ where $mutex.monitor = monitor$. No-reply with **(leave-mutex mutex state)**.

(enter-mutex resource from waiters state) **procedure**

returns: updated state

The **enter-mutex** procedure calls **(gen-server:reply from 'ok)** to reply to the caller waiting to enter the mutex. It adds a **<mutex>** tuple with $resource = resource$, $process = from.process$, $monitor = (monitor\ process)$, $count = 1$, and $waiters = waiters$ to $state$.

(leave-mutex mutex state) **procedure**

returns: updated state

The **leave-mutex** procedure calls **(demonitor&flush mutex.monitor)**. If $mutex.waiters = ()$, it returns **(remq mutex state)**. Otherwise, it returns **(enter-mutex mutex.resource (car mutex.waiters) (cdr mutex.waiters) (remq mutex state))**.

(deadlock? process mutex state) **procedure**

returns: a boolean

The **deadlock?** procedure returns **#t** if $process$ would deadlock waiting for $mutex$. Let $owner = mutex.process$. If $owner = process$, return **#t**. Otherwise, find the mutex $waiting \in state$ where **#(owner _)** $\in waiting.waiters$. If no such $waiting$ exists, return **#f**. Otherwise, return **(deadlock? process waiting state)**.

6.3 Programming Interface

<code>(gatekeeper:start&link)</code>	procedure
<code>returns: #(ok <i>pid</i>) #(error <i>reason</i>)</code>	

The `gatekeeper:start&link` procedure calls `(gen-server:start&link 'gatekeeper)`.

<code>(gatekeeper:enter <i>resource timeout</i>)</code>	procedure
<code>returns: ok</code>	

The `gatekeeper:enter` procedure calls `(gen-server:call 'gatekeeper #(enter resource) timeout)` to enter the mutex for *resource*. If it returns *e* \neq `ok`, it raises exception *e*.

<code>(gatekeeper:leave <i>resource</i>)</code>	procedure
<code>returns: ok</code>	

The `gatekeeper:leave` procedure calls `(gen-server:call 'gatekeeper #(leave resource))` to leave the mutex for *resource*. If it returns *e* \neq `ok`, it raises exception *e*.

<code>(with-gatekeeper-mutex <i>resource timeout body₁ body₂ ...</i>)</code>	syntax
<code>expands to: (\$with-gatekeeper-mutex '<i>resource timeout</i> (lambda () <i>body₁ body₂ ...</i>))</code>	

The `with-gatekeeper-mutex` form executes the body expressions in a dynamic context where the calling process owns *resource*, which must be an identifier. The *timeout* expression specifies how long the caller is willing to wait to enter the mutex for *resource* as defined by `gen-server:call`. The internal `$with-gatekeeper-mutex` procedure is defined as follows:

```
(define ($with-gatekeeper-mutex resource timeout body)
  (dynamic-wind
    (lambda () (gatekeeper:enter resource timeout))
    body
    (lambda () (gatekeeper:leave resource))))
```


Chapter 7

Supervisor

7.1 Introduction

In a fault tolerant system, faults must first be observed and then acted upon. A `supervisor` monitors child processes for failure and can be composed into a hierarchy to monitor for faults within other supervisors.

The principles of supervisors and supervision hierarchies can be found in Joe Armstrong's thesis [1] or *Programming Erlang—Software for a Concurrent World* [2]. Documentation for Erlang's `supervisor` is available online [12]. Source code for the Erlang Open Telecom Platform can be found online [8]. The source code for `supervisor` is part of `stdlib` and can be found in `/lib/stdlib/src/supervisor.erl`.

Patterns for Fault Tolerant Software [15] is a good reference for understanding the mindset of creating fault tolerant systems.

7.2 Theory of Operation

A *supervisor* is a gen-server which is responsible for starting, stopping, and monitoring its child processes. A supervisor observes its children, and when a failure occurs, restarts child processes.

A *watcher* is a supervisor which is configured to only observe the children. A watcher interface is provided for convenience.

A supervisor can be configured to restart individual children when those children fail, or to restart all children when any child fails. This is called the restart *strategy*. A strategy of `one-for-one` indicates that when a child process terminates, it should be restarted; only that child process is affected. A strategy of `one-for-all` indicates that when a child process terminates and should be restarted, all other child process are terminated and then restarted.

A supervisor maintains a list of times of when a restart occurs. When a child fails and is to be restarted, a timestamp is added to the *restarts* list. A maximum restart frequency is represented as an *intensity* and a *period* of time. If more than *intensity* restarts occur in a *period* of time, the supervisor terminates all child processes and then itself. This prevents the possibility of an infinite

cycle of child process termination and restarts.

A supervisor is started with a list of child specifications. These specifications are used to start child processes from within the supervisor process during initialization.

Child specifications can be added to a supervisor at run time. These dynamic children will not be automatically restarted if the supervisor itself terminates and is restarted.

state (define-state-tuple <supervisor-state> strategy intensity period children restarts)

- **strategy** defines how the supervisor processes a child termination: **one-for-one** or **one-for-all**.
- **intensity** is the maximum restart intensity.
- **period** is the maximum restart period in milliseconds.
- **children** is a list of <child> tuples with the most recently started child first.
- **restarts** is an ordered list of times when restarts have occurred.

(define-tuple <child> pid name thunk restart-type shutdown type)

pid stores the child process or **#f**. The remaining fields are copied from the child specification described below.

init The **init** procedure validates the startup arguments and starts the initial child processes. Invalid startup arguments cause the supervisor to fail to start. If any child fails to start, all started children are terminated and the supervisor fails to start.

This process traps exits so that it can detect child exits, as well as the **EXIT** message from the parent process.

An invalid argument results in a specific error reason that includes the invalid input.

- **invalid-strategy** *strategy*
- **invalid-intensity** *intensity*
- **invalid-period** *period*

An invalid child specification during initialization will result in **(error (start-specs reason))** where *reason* is one of the reasons listed in the programming interface below.

terminate The **terminate** procedure shuts each child process down in order (most recently added first).

handle-call The **handle-call** procedure processes the following messages:

- **(start-child child-spec)**: Validates the *child-spec*, starts the child, adds it to the state, and replies with **(ok pid)** where *pid* is the new child process.

If a child specification of the same name already exists, `#(error already-present)` is returned. If the child process was already started `#(error #(already-started pid))` is returned.

A successfully started child is linked to the supervisor, an event is fired to the event manager to log the start, and `#(ok pid)` is returned. If the *pid* already occurs in the children list, then `start-child` returns `#(error #(duplicate-process pid))`.

If the child process start function returns `ignore`, the child specification is added to the supervisor, and the function returns `#(ok #f)`.

If the child process start function returns `#(error reason)`, then `start-child` returns `#(error reason)`.

If the child process start function exits with *reason*, `#(error reason)` is returned.

If the child process start function returns *other* values `#(error #(bad-return-value other))` is returned.

- `#(restart-child name)`: Finds a child by *name*, verifies that it is not currently running, then starts that child.

If the child process is already running, `#(error running)` is returned. If the child specification does not exist, `#(error not-found)` is returned.

A successfully started child is linked to the supervisor, an event is fired to the event manager to log the start, and `#(ok pid)` is returned. If the *pid* already occurs in the children list, then `restart-child` returns `#(error #(duplicate-process pid))`.

If the child process start function returns `ignore`, the child specification is added to the supervisor and the function returns `#(ok #f)`.

If the child process start function returns `#(error reason)`, then `restart-child` returns `#(error reason)`.

If the child process start function exits with *reason*, `#(error reason)` is returned.

If the child process start function returns *other* values `#(error #(bad-return-value other))` is returned.

- `#(delete-child name)`: Finds a child by *name*, verifies that it is not currently running, then removes the child specification from the state and returns `ok`.

If the child process is running, `#(error running)` is returned. If the child specification does not exist, `#(error not-found)` is returned.

- `#(terminate-child name)`: Finds a child by *name* and terminates it if it is running. The child *pid* is updated to `#f` and returns `ok`.

If the child specification does not exist, `#(error not-found)` is returned.

- `get-children`: Returns the state's `children` field.

`handle-cast` The `handle-cast` procedure does not process any messages.

handle-info The **handle-info** procedure processes the following message:

- **#(EXIT *pid* *reason*)**: Find *pid* in the children list and apply the restart strategy. An unknown *pid* is ignored.

When the child specification *restart-type* is **permanent** or **transient** the current timestamp is prepended to the *restarts* list. The list is then pruned based on the *period*. If the resulting list length \leq *intensity*, the supervisor continues. Otherwise, the supervisor terminates with reason **shutdown**.

Internally, the (**shutdown** *pid* *x*) function kills child processes and returns the exit reason. This function is used by **terminate**, **terminate-child**, and during a failed **init**. The following steps are necessary to defend against a “naughty” child which unlinks from the supervisor.

- Monitor *pid* to protect against a child process which may have called **unlink**.
- Unlink *pid* to stop receiving **EXIT** messages from *pid*.
- An **EXIT** message may already exist for *pid*. If it does, then wait for the **DOWN** message, and return the exit reason.
- If *x* = **brutal-kill**, kill *pid* with reason **kill** and wait for the **DOWN** message to determine the exit reason.
- Otherwise, *x* is a *timeout*. kill *pid* with reason **shutdown** and wait for the **DOWN** message to determine the exit reason. If a *timeout* occurs, kill *pid* with reason **kill**, and wait for the **DOWN** message to determine the exit reason.

7.3 Design Decisions

Our initial implementation did not automatically link to child processes, but this led to unexpected behavior when child processes neglected to link to the supervisor. Therefore, this implementation links to all child processes.

7.4 Programming Interface

supervisor:start&link and **supervisor:start-child** use a child specification. A child specification is defined as:

child-spec \rightarrow **#(name** *thunk* *restart-type* *shutdown* *type*)

name is a symbol unique to the children within the supervisor.

thunk is a procedure that should spawn a process and link to the supervisor process, then return **#(ok** *pid*) or **#(error** *reason*) or **ignore**. Typically, the *thunk* will call **gen-server:start&link** which provides the appropriate behavior and return value.

restart-type is a symbol with the following meaning:

- A **permanent** child process is always restarted.
- A **temporary** child process is never restarted.

- A **transient** child process is only restarted if it terminates with an exit reason other than **normal** or **shutdown**.
- A **watch-only** child process is never restarted, and its child specification is removed from the supervisor when it terminates.

shutdown defines how a child process should be terminated.

- **brutal-kill** indicates that the child process will be terminated using (**kill** *pid* **kill**).
- A fixnum > 0 represents a timeout. The supervisor will use (**kill** *pid* **shutdown**) and wait for an exit signal. If no exit signal is received within the timeout, the child process will be terminated using (**kill** *pid* **kill**). **infinity** can be used if and only if the *type* of the process is **supervisor**.

The *type* is useful for validating the *shutdown* parameter, but is otherwise unused. It may be useful in conjunction with **supervisor:get-children** to generate a tree of the running supervision hierarchy.

```
type → supervisor
      | worker
```

Invalid child specifications will result in specific error reasons which include the invalid input.

- **#(invalid-name** *name*)
- **#(invalid-thunk** *thunk*)
- **#(invalid-restart-type** *restart-type*)
- **#(invalid-type** *type*)
- **#(invalid-shutdown** *shutdown*)
- **#(invalid-child-spec** *spec*)

(**supervisor:start&link** *name strategy intensity period start-specs*) **procedure**
returns: **#(ok** *pid*) **| #(error** *reason*)

The **supervisor:start&link** procedure creates a new **supervisor** gen-server using **gen-server:start&link**. *name* is the registered name of the process. For an anonymous server, **#f** may be specified.

```
strategy → one-for-one
           | one-for-all
```

```
intensity → a fixnum >= 0
```

```
period → a fixnum > 0
```

```
start-specs → (child-spec ...)
```

(**supervisor:start-child** *supervisor child-spec*) **procedure**

returns: `#(ok pid) | #(error reason)`

This procedure dynamically adds the given *child-spec* to the *supervisor* which starts a child process.

The `supervisor:start-child` procedure calls `(gen-server:call supervisor #(start-child child-spec) infinity)`.

(supervisor:restart-child <i>supervisor</i> <i>name</i>)	procedure
returns: <code> #(ok <i>pid</i>) #(error <i>reason</i>)</code>	

This procedure restarts a child process identified by *name*. The child specification must exist, and the child process must not be running.

The `supervisor:restart-child` procedure calls `(gen-server:call supervisor #(restart-child name) infinity)`.

(supervisor:delete-child <i>supervisor</i> <i>name</i>)	procedure
returns: <code> ok #(error <i>reason</i>)</code>	

This procedure deletes the child specification identified by *name*. The child process must not be running.

The `supervisor:delete-child` procedure calls `(gen-server:call supervisor #(delete-child name) infinity)`.

(supervisor:terminate-child <i>supervisor</i> <i>name</i>)	procedure
returns: <code> ok #(error <i>reason</i>)</code>	

This procedure terminates the child process identified by *name*. The child specification must exist, but the child process does not need be running.

The `supervisor:terminate-child` procedure calls `(gen-server:call supervisor #(terminate-child name) infinity)`.

(supervisor:get-children <i>supervisor</i>)	procedure
returns: a list of <code><child></code> tuples	

This procedure returns the *supervisor* internal representation of child specifications.

The `supervisor:get-children` procedure calls `(gen-server:call supervisor get-children infinity)`.

7.5 Published Events

A supervisor can notify the event manager of the same events as a gen-server, as well as the following events.

```
event → <supervisor-error>
      | <child-start>
      | <child-end>
```

<supervisor-error>	event
---------------------------------	--------------

timestamp: the time the event occurred
supervisor: the supervisor's process id
error-context: the context in which the event occurred
reason: the reason for the error
child-pid: the child's process id
child-name: the child's name

This event is fired when the supervisor fails to start its children, fails to restart its children, or when it has exceeded the maximum restart frequency.

<child-start>	event
----------------------------	--------------

timestamp: the time the event occurred
supervisor: the supervisor's process id
pid: the child's process id
name: the child's name
restart-type: the child's restart-type
shutdown: the child's shutdown
type: the child's type

This event is fired after the child start procedure has returned a valid value.

<child-end>	event
--------------------------	--------------

timestamp: the time the event occurred
pid: the child's process id
killed: 1 indicates the supervisor terminated the child, 0 otherwise
reason: the reason the child has terminated

This event is fired after the supervisor terminates a child process, and after the supervisor detects a failure in a child.

7.6 Watcher Interface

(watcher:start&link <i>name</i>)	procedure
---	------------------

returns: #(ok *pid*) | #(error *reason*)

The `watcher:start&link` procedure creates a supervisor with a strategy of `one-for-one`, an intensity of 0, a period of 1, and no children.

name is the registered name of the process. For an anonymous server, `#f` may be specified.

(watcher:start-child <i>watcher name shutdown thunk</i>)	procedure
---	------------------

returns: #(ok *pid*) | #(error *reason*)

The `watcher:start-child` procedure calls `(supervisor:start-child watcher #(name thunk watch-only shutdown worker))`.

(watcher:shutdown-children <i>watcher</i>)	procedure
---	------------------

returns: ok

The `watcher:shutdown-children` procedure terminates and deletes each `watch-only` child in *watcher*.

Chapter 8

Application

8.1 Introduction

The application is a single gen-server named `application` that manages the lifetime of the program. It links to a process, typically the root supervisor, and shuts down the program when requested by `application:shutdown` or when the linked process dies.

8.2 Theory of Operation

state The application state is the process returned by the *starter* of `application:start`. It is typically the root supervisor. We refer to this variable as *process*. It may also be `#f` after `handle-info` receives the exit message for the process.

init The application `init` procedure takes a *starter* procedure. It calls (*starter*) and checks the return value *r*. If *r* = `#(ok process)`, it links to *process*, traps exits so that it receives exit messages from *process* and `application:shutdown`, and returns `#(ok process)`. If *r* = `#(error reason)`, it returns `#(stop reason)`.

terminate The application `terminate` procedure shuts down *process*. When *process* is not `#f`, it kills *process* with reason `shutdown` and waits indefinitely for it to terminate. Then it calls (`exit-process exit-code`), where *exit-code* is initially 2 but set to the value passed to `application:shutdown`. In this way, the exit code can be used to determine if the application shut down normally.

handle-call The application `handle-call` procedure raises an exception on all messages.

handle-cast The application `handle-cast` procedure raises an exception on all messages.

handle-info The application `handle-info` procedure handles the following message:

- `#(EXIT p reason)`: If $p = process$, return `#(stop reason #f)`. Otherwise, return `#(stop reason process)`.

<code>(exit-process <i>exit-code</i>)</code>	procedure
returns: never	

The `exit-process` procedure flushes the console output port, ignoring any exceptions, and then calls `(osi_exit exit-code)`.

8.3 Programming Interface

<code>(application:start <i>starter</i>)</code>	procedure
returns: ok	

The `application:start` procedure calls `(gen-server:start 'application starter)`. If it returns `#(ok _)`, `application:start` returns ok. If it returns `#(error reason)`, `application:start` calls `(console-event-handler #(application-start-failed reason))` and `(exit-process 1)`.

<code>(application:shutdown [<i>exit-code</i>])</code>	procedure
returns: unspecified	

The `application:shutdown` procedure kills the `application` process with reason `shutdown`. The *exit-code* defaults to 0, indicating normal shutdown. The procedure does not wait for the `application` process to terminate so that it can be called from a process managed by the supervision hierarchy without causing a deadlock on shutdown. If the `application` process does not exist, `application:shutdown` calls `(exit-process exit-code)`.

Chapter 9

Database Interface

9.1 Introduction

The database (`db`) interface is a gen-server which provides a basic transaction framework to retrieve and store data in a SQLite database. It provides functions to use transactions (directly and lazily).

The low-level SQLite interface can be found in the operating system interface design (see Chapter 2).

Other SQLite resources are available online [23] or in The Definitive Guide to SQLite [20].

9.2 Theory of Operation

The `db` gen-server serializes internal requests to the database. For storage and retrieval of data, each transaction is processed in turn by a separate linked process. The gen-server does not block waiting for this process to finish so that it can maintain linear performance by keeping its inbox short. The return value of the transaction is returned to the caller or an error is generated without tearing down the gen-server.

To facilitate logging, the `db` gen-server can lazily open a transaction. In order to allow other processes access to the database, lazy transactions should be closed occasionally. To support this, it tracks a count of entries in the current transaction. A transaction is committed when the threshold of 10,000 is reached, the message queue of the `db` is empty, or when a direct transaction is requested. Each database is created with write-ahead logging enabled to prevent write operations from blocking on queries made from another connection.

SQLite has three types of transactions: deferred, immediate, and exclusive. This interface uses only immediate transactions to simplify the handling of the `SQLITE_BUSY` error. Using immediate transactions means that `SQLITE_BUSY` will only occur during `BEGIN IMMEDIATE`, `BEGIN TRANSACTION`, `COMMIT`, and `ROLLBACK`¹ statements. For each of these statements, when a `SQLITE_BUSY` occurs, the code waits for a brief time, then retries the statement. The wait times in milliseconds follow the pattern (2 3 6 11 16 21 26 26 26 51 51 . #0=(101 . #0#)), and up to 500 retries are attempted before exiting with `#{db-retry-failed sql count}`. When the retry count is positive,

¹Our testing showed that `ROLLBACK` returns `SQLITE_BUSY` only when a `COMMIT` for the same transaction returned `SQLITE_BUSY`. This framework never causes that situation to occur, but it guards against it anyway.

it is logged to the event manager along with the total duration with a `<transaction-retry>` event.

<code><transaction-retry></code>	event
<i>timestamp</i> : timestamp from <code>erlang:now</code>	
<i>database</i> : database filename	
<i>duration</i> : duration in milliseconds	
<i>count</i> : retry count	
<i>sql</i> : query	

The `db` gen-server uses the operating system interface to interact with SQLite. To prevent memory leaks, raw database handles are wrapped in a Scheme record and registered with a guardian.

```
state (define-state-tuple <db-state> filename db cache queue worker)
```

- `filename` is the database specified when the server was started.
- `db` is the database record.
- `cache` is a hash table mapping SQL strings to SQLite prepared statements.
- `queue` is a queue of log and transaction requests.
- `worker` is the pid of the active worker or `#f`.

dictionary parameters

- `current-database` stores a Scheme record:

```
(define-record-type database (fields (mutable handle)))
```

The `handle` is set to `#f` when the database is closed.
- `statement-cache` stores a Scheme record:

```
(define-record-type cache (fields (immutable ht) (mutable waketime) (mutable lazy-statements)))
```

The `waketime` is the next time the cache will attempt to remove dead entries.

The hash table, `ht`, maps SQL strings to a Scheme record:

```
(define-record-type entry (fields (immutable stmt) (mutable timestamp)))
```

When a SQL string is not found in the cache, `PrepareStatement` is used with the `current-database` to make a SQLite statement. The raw statement handle is stored in a Scheme record:

```
(define-record-type statement (fields (immutable handle) (immutable database)))
```

The `statement` record is not registered with a guardian. The statement is finalized using `FinalizeStatement` when it is removed from the cache. `CloseDatabase` will finalize any remaining statements associated with the database.

When a SQL string is found in the cache, the entry's `timestamp` is updated. Entries older than 5 minutes will be removed from the cache.

Accessing the cache may exit with reason `reason #(db-error prepare error sql)`, where *error* is a SQLite error pair.

The `lazy-statements` list contains `statement` records created by `lazy-execute`. These statements are finalized when a transaction completes.

init The `init` procedure takes a filename and mode symbol and attempts to open that database, setting `journal_mode` to “wal” if *mode* is `create`. The handle returned from `OpenDatabase` is wrapped in a *database* record that is registered with a guardian. The garbage collector is hooked so that dead databases are closed even if the `db` gen-server fails to close them for any reason.

The gen-server traps exits so that it can close the database in its `terminate` procedure.

terminate The `terminate` flushes the queue and closes the database.

handle-call The `handle-call` procedure processes the following messages:

- `#(transaction f)`: Add this transaction along with the *from* argument to `handle-call` to the queue. Process the queue.
- `filename`: Return the database filename.
- `stop`: Flush the queue and stop with reason `normal`, returning `stopped` to the caller.

handle-cast The `handle-cast` procedure processes the following message:

- `#(log sql bindings)`: Add this tuple to the queue. Process the queue.

handle-info The `handle-info` procedure processes the following messages:

- `timeout`: Remove old entries from the statement cache.
- `#(EXIT worker-pid normal)`: The worker finished the previous request successfully. Process the queue.
- `#(EXIT worker-pid reason)`: The worker failed to process the previous request. Flush the queue and stop with *reason*.

9.3 Design Decisions

There is a one-to-one relationship between a SQLite database handle and the `db` gen-server. For clarity, the database handle and a SQLite statement cache are implemented in terms of Erlang process dictionary parameters.

An alternate approach for logging was already explored where a transaction was not lazily opened. Such an approach means that when a third party tool tries to access the database, it will hang until the transaction is complete.

A commit threshold of 10,000 was chosen because it was large enough to minimize the cost of a transaction but small enough to execute simple queries in less than one second.

9.4 Programming Interface

(db:start&link *name filename mode*) **procedure**
returns: #(ok *pid*) | #(error *error*)

The `db:start&link` procedure creates a new `db` gen-server using `gen-server:start&link`.

name is the registered name of the process. For an anonymous server, `#f` may be specified.

filename is the path to a SQLite database.

mode is one of the following symbols used to pass SQLite flags to `OpenDatabase`:

- `read-only` uses the SQLite flag `SQLITE_OPEN_READONLY`.
- `open` uses the SQLite flag `SQLITE_OPEN_READWRITE`.
- `create` combines the SQLite flags (`logor SQLITE_OPEN_READWRITE SQLITE_OPEN_CREATE`).

The SQLite constants can be found in `sqlite3.h` or online [23].

This procedure may return an *error* of `#(db-error open error filename)`, where *error* is a SQLite error pair.

(db:stop *who*) **procedure**
returns: `stopped`

The `db:stop` procedure calls `(gen-server:call who stop infinity)`.

(with-db [*db filename flags*] *body₁ body₂ ...*) **syntax**
expands to:
`(let ([db (sqlite:open filename flags)])
 (on-exit (sqlite:close db)
 body1 body2 ...))`

The `with-db` macro opens the database in *filename*, executes the statements in the body, and closes the database before exiting. This is a suitable alternative to starting a `gen-server` when you need to query a database using a separate SQLite connection, and you do not need to cache prepared SQL statements.

(db:filename *who*) **procedure**
returns: the database filename

The `db:filename` procedure calls `(gen-server:call who filename)`.

(db:log *who sql . bindings*) **procedure**
returns: `ok`

The `db:log` procedure calls `(gen-server:cast who #(log sql bindings))`. *sql* is a SQL string, and *bindings* is a list of values to be bound in the query. Because `db:log` does not wait for a reply from the server, any error in processing the request will crash the server.

<code>(db:transaction who f)</code>	procedure
returns: <code> #(ok result) #(error error)</code>	

The `db:transaction` procedure calls `(gen-server:call who #(transaction f) infinity)`.

f is a thunk which returns a single value, *result*. `execute`, `lazy-execute`, and `columns` can be used inside the procedure *f*.

result is the successful return value of *f*. Typically, this is a list of rows as returned by a `SELECT` query.

error is the failure reason of *f*.

<code>(transaction db body ...)</code>	syntax
expands to:	

```
(match (db:transaction db (lambda () body ...))
  [#(ok ,result) result]
  [#(error ,reason) (raise reason)])
```

The `transaction` macro runs the body in a transaction and returns the result when successful and exits when unsuccessful.

<code>(execute sql . bindings)</code>	procedure
returns: a list of rows where each row is a vector of data in column order as specified in the <i>sql</i> statement	

`execute` should only be used from within a thunk *f* provided to `db:transaction`.

sql is mapped to a SQLite statement using the `statement-cache`. The *bindings* are then applied using `BindStatement`. The statement is then executed using `StepStatement`. The results are accumulated as a list, and the statement is reset using `ResetStatement` to prevent the statement from locking parts of the database.

This procedure may exit with reason `#(db-error prepare error sql)`, where *error* is a SQLite error pair.

<code>(lazy-execute sql . bindings)</code>	procedure
returns: a thunk	

`lazy-execute` should only be used from within a thunk *f* provided to `db:transaction`.

A new SQLite statement is created from *sql* using `PrepareStatement` so that the statement won't interfere with any other queries. The statement is added to the `lazy-statements` list of the `statement-cache` and is finalized when the transaction completes. The *bindings* are then applied using `BindStatement`. A thunk is returned which, when called, executes the statement using `StepStatement`. The thunk returns one row of data or `#f`.

This procedure may exit with reason `#(db-error prepare error sql)`, where *error* is a SQLite error pair.

<code>(execute-sql db sql . bindings)</code>	procedure
--	------------------

returns: a list of rows where each row is a vector of data in column order as specified in the *sql* statement

execute-sql should only be used for statements that do not need to be inside a transaction, such as a one-time query.

sql is prepared into a SQLite statement for use with *db*, executed via `sqlite:execute` with the specified *bindings*, and finalized.

This procedure may exit with reason `#(db-error prepare error sql)`, where *error* is a SQLite error pair.

(columns *sql*) **procedure**

returns: a vector of column names in order as specified in the *sql* statement

columns should only be used from within a thunk *f* provided to `db:transaction`.

sql is mapped to a SQLite statement using the **statement-cache**. The statement columns are then retrieved using `GetStatementColumns`.

(parse-sql *x* [*symbol*->*sql*]) **procedure**

returns: two values: a query string and a list of syntax objects for the arguments

The **parse-sql** procedure is used by macro transformers to take syntax object *x* and produce a query string and associated arguments according to the patterns below. When one of these patterns is matched, the *symbol*->*sql* procedure is applied to the remaining symbols of the input before they are spliced into the query string, as if by `(format "~a" (symbol->sql sym))`. By default, *symbol*->*sql* is the identity function.

- **(insert *table* ([*column* *e*₁ *e*₂ ...] ...) ...)**

The **insert** form generates a SQL insert statement. The *table* and *column* patterns are SQL identifiers. Any *e* expression that is `(unquote exp)` is converted to ? in the query, and *exp* is added to the list of arguments. All other expressions are spliced into the query string.

- **(update *table* ([*column* *e*₁ *e*₂ ...] ...) *where* ...)**

The **update** form generates a SQL update statement. The *table* and *column* patterns are SQL identifiers. Any *e* or *where* expression that is `(unquote exp)` is converted to ? in the query, and *exp* is added to the list of arguments. All other expressions are spliced into the query string.

- **(delete *table* *where* ...)**

The **delete** form generates a SQL delete statement. The *table* pattern is a SQL identifier. Any *where* expression that is `(unquote exp)` is converted to ? in the query, and *exp* is added to the list of arguments. All other expressions are spliced into the query string.

(sqlite:bind *stmt* *bindings*) **procedure**

returns: unspecified

The **sqlite:bind** procedure binds the variables in statement record instance *stmt* with the list of *bindings*. It resets the statement before binding the variables.

(sqlite:close *db*) **procedure**

returns: unspecified

The `sqlite:close` procedure closes the database associated with database record instance *db*.

(sqlite:columns *stmt*) **procedure**

returns: a vector of column names

The `sqlite:columns` procedure returns a vector of column names for the statement record instance *stmt*.

(sqlite:execute *stmt bindings*) **procedure**

returns: a list of rows where each row is a vector of data in column order

The `sqlite:execute` procedure calls `(sqlite:bind stmt bindings)` to bind any variables and then iteratively calls `(sqlite:step stmt)` to build the resulting list of rows. It resets the statement when the procedure exits.

(sqlite:expanded-sql *stmt*) **procedure**

returns: a string

The `sqlite:expanded-sql` procedure returns the SQL string expanded with the binding values for the statement record instance *stmt*.

(sqlite:finalize *stmt*) **procedure**

returns: unspecified

The `sqlite:finalize` procedure finalizes the statement record instance *stmt*.

(sqlite:open *filename flags*) **procedure**

returns: a database record instance

The `sqlite:open` procedure opens the SQLite database in file *filename* with *flags* specified by `sqlite3_open_v2` [23]. The constants `SQLITE_OPEN_CREATE`, `SQLITE_OPEN_READONLY`, and `SQLITE_OPEN_READWRITE` are exported from the `(swish db)` library.

(sqlite:prepare *db sql*) **procedure**

returns: a statement record instance

The `sqlite:prepare` procedure returns a statement record instance for the *sql* statement in the database record instance *db*.

(sqlite:sql *stmt*) **procedure**

returns: a string

The `sqlite:sql` procedure returns the unexpanded SQL string for the statement record instance *stmt*.

(sqlite:step *stmt*) **procedure**

returns: a vector of data in column order or **#f**

The `sqlite:step` procedure steps the statement record instance *stmt* and returns the next row vector in column order or **#f** if there are no more rows.

Chapter 10

Log Database

10.1 Introduction

The log database is a single gen-server named `log-db` that uses the database interface (see Chapter 9) to log system events (see Chapter 5).

10.2 Theory of Operation

10.2.1 Initialization

The `log-db` gen-server handles startup and setup through two separate procedures. Startup uses the `db:start&link` procedure to connect to the the SQLite log database specified by the (`log-file`) parameter. It creates the file if it does not exist, but otherwise startup does not modify the database.

Setup makes sure the schema of the log database has been created and is up-to-date. A unique symbol identifying the schema version is stored in a table named `version`. This allows the software to upgrade between known schema versions and to exit with an error when it encounters an unsupported database version. These schema updates happen within a database transaction so that if there is an error, the changes are rolled back.

Setup calls `event-mgr:set-log-handler` after updating the schema. This registers the `log-db` to log system events. It also calls `event-mgr:flush-buffer`. This causes the event manager to stop buffering startup events and the `log-db` to log the events that were buffered.

Finally, setup sends a `<system-attributes>` event so that `log-db` receives and logs it.

Once the `log-db` gen-server has been setup, it continues to receive events from the system event manager. It converts events that it recognizes into insertions to the log database. Events that it does not recognize are ignored.

The tables are pruned using insert triggers to hold 90 days of information. To keep the insert operations fast, the timestamp columns are indexed, and the pruning deletes no more than 10 rows per insert.

10.2.2 Extensions

An application typically produces events beyond those that are part of Swish and may wish to log them in the same log database file where the Swish events are logged. The `log-db` design allows for this type of extension.

The `log-db:setup` procedure takes a list of `<event-logger>` tuples. Each logger represents an extension to the log database schema and contains two procedures, `setup` and `log`. The `log-db:setup` procedure calls the `setup` procedure of logger to make sure that its portion of the schema has been created and is up-to-date. Then, when `log-db` receives an event, it calls the `log` procedure of each logger. If the event is recognized by that portion of the schema, the `log` procedure inserts or updates data in the log database. Otherwise, the procedure ignores that event.

Additionally, the version table does not store a single schema version. Instead, it stores schema versions associated with names. The `setup` procedure of an `<event-logger>` uses an unique name for its portion of the schema and the `log-db:version` procedure to retrieve and set its version.

The schema and logging for Swish events is implemented as an `<event-logger>` defined by `swish-event-logger` and using the schema version name `swish`. An application that wishes to use this logging must provide `swish-event-logger` in the list to `log-db:setup`. If the application wishes to log Swish events in a different structure, it can omit the `swish-event-logger` and provide its own logger with its own schema. However, doing so makes the application more brittle with respect to changes in the Swish implementation.

10.3 Programming Interface

<code><event-logger></code>	tuple
-----------------------------------	--------------

setup: procedure of no arguments that makes sure this portion of the schema is created and up-to-date

log: procedure of one argument, an event, that logs the event if it recognizes it and otherwise ignores it

<code>(log-db:start&link)</code>	procedure
--------------------------------------	------------------

returns: `#(ok pid) | #(error error)`

The `log-db:start&link` procedure creates a new `db` gen-server named `log-db` using `db:start&link`. It uses the value of the `(log-file)` parameter as the path to the SQLite database and specifies `create` mode.

<code>(log-db:setup loggers)</code>	procedure
-------------------------------------	------------------

returns: `ignore | #(error error)`

The argument *loggers* is a list of `<event-logger>` tuples. The `log-db:setup` makes sure the `log-db` is setup to run by doing the following in order.

1. Initialize or upgrade the database schema from within a `db:transaction` call. It does this by calling the `setup` procedure of each logger.

2. Register a procedure with `event-mgr:set-log-handler` to have the `log-db` gen-server log events it recognizes. When this procedure receives an event, it calls the `log` procedure of each logger.
3. Call `event-mgr:flush-buffer` to stop buffering system events and apply the log handler to the events already buffered.
4. Send a `<system-attributes>` event.

If everything succeeds, the procedure returns `ignore`. If either the `db:transaction` or `event-mgr:set-log-handler` indicate an error, the procedure returns that error.

(log-db:version *name* [*version*]) **procedure**

name: symbol identifying the schema

version: string specifying the version of the schema

When called with one argument, `log-db:version` retrieves the version associated with *name* from the database and returns it as a string. It returns `#f` if no version associated with *name* is stored in the database.

When called with two arguments, it stores *version* as the version associated with *name* in the database.

(log-db:get-instance-id) **procedure**

returns: a string

`log-db:setup` associates a globally unique identifier with the database file. The `log-db:get-instance-id` function caches and returns that identifier.

swish-event-logger **property**

The `swish-event-logger` is an `<event-logger>` tuple that defines the schema for Swish events. It uses the name `swish` to store its schema version.

**(create-table *name*
 (*field type . inline*)
 ...)** **syntax**

expands to: (execute "create table if not exists ...")

The `create-table` syntax describes the schema of a single table and expands into a call to `execute` to create the table if no table with that name already exists. The name of the table, *name*, and of each field, *field*, are converted from Scheme to SQL identifiers by replacing hyphen characters with underscores and eliminating any non-alphanumeric and non-underscore characters. The SQL definition of each field is produced by joining the converted field name, the *type* and any additional *inline* arguments into a space separated string.

**(define-simple-events *create handle*
 (*name clause* ...)
 ...)** **syntax**

expands to: A definition of the *create* and *handle* procedures

The **define-simple-events** syntax is used to log tuple types by inserting a row into a table with the same name and the same fields. Each *name* is a tuple type. Each *clause* is a valid **create-table** clause for one of the fields in that tuple type.

It defines *create* as a procedure of 0 arguments that consists of a (**create-table** *name clause* ...) for each tuple in the **define-simple-events**. This means that the name of the tuple type and each field are converted to SQL names by the **create-table** syntax.

It defines *handle* as a procedure of 1 argument, an event. If the event is one of the tuple types in the **define-simple-events**, it calls **db:log** with an insert statement applying **coerce** to each value. If the event is unrecognized, it returns **#f**.

(coerce <i>x</i>)	procedure
returns: a Scheme object	

The argument *x* is a Scheme object mapped to a SQLite value.

<i>type</i>	transformation
<i>string</i>	<i>string</i>
<i>bytevector</i>	<i>bytevector</i>
<i>number</i>	<i>number</i> , if it fits in 64 bits
<i>symbol</i>	symbol->string
<i>date</i>	format-rfc2822
<i>process</i>	integer key for <i>process</i> , unique in this database instance
<i>condition</i>	a string containing #(error reason) where the <i>reason</i> is obtained from display-condition
<i>continuation-condition</i>	a string containing #(error reason stack) where the <i>stack</i> is obtained from dump-stack

coerce passes **#f** through unmodified which SQLite interprets as NULL. Other values are converted to string using **write**.

10.4 Published Events

<system-attributes>	event
<i>timestamp</i> : timestamp from erlang:now	
<i>date</i> : date from current-date	
<i>software-version</i> : software version string	
<i>computer-name</i> : computer name from osi::GetComputerName	

The **<system-attributes>** event is sent exactly once, when **log-db:setup** is called.

Chapter 11

System Statistics

11.1 Introduction

The system uses a single gen-server named `statistics` to periodically query statistics about the system, such as memory usage.

11.2 Theory of Operation

When the `statistics` gen-server starts, it posts a `<statistics>` event with `reason = startup`. Every five minutes thereafter, it posts a `<statistics>` event with `reason = update`. If the computer sleeps or hibernates, the gen-server posts a `<statistics>` event with `reason = suspend`. When the computer awakens, the gen-server posts a `<statistics>` event with `reason = resume`. When the gen-server terminates, it posts a `<statistics>` event with `reason = shutdown`.

The `<statistics>` event is handled by the `log-db` gen-server (see Chapter 10), which adds the data to the statistics table in the log database.

11.3 Programming Interface

```
(statistics:start&link) procedure  
returns: #(ok pid) | #(error error)
```

The `statistics:start&link` procedure creates a new gen-server named `statistics` using `gen-server:start&link`. It then posts a `<statistics>` event with `reason = startup`.

```
(statistics:resume) procedure
```

The `statistics:resume` procedure casts a message to the `statistics` gen-server that causes it to publish a `<statistics>` event with `reason = resume`. This procedure is aliased to `app:resume` and called from the operating system interface.

```
(statistics:suspend) procedure
```

The `statistics:suspend` procedure casts a message to the `statistics` gen-server that causes it to publish a `<statistics>` event with `reason = suspend`. This procedure is aliased to `app:suspend` and called from the operating system interface.

11.4 Published Events

<code><statistics></code>	event
<i>timestamp</i> : timestamp from <code>erlang:now</code>	
<i>date</i> : date from <code>current-date</code>	
<i>reason</i> : <code>startup</code> , <code>update</code> , <code>suspend</code> , <code>resume</code> , or <code>shutdown</code>	
<i>bytes-allocated</i> : Scheme heap size from <code>bytes-allocated</code>	
<i>osi-bytes-used</i> : C heap size from <code>osi_get_bytes_used</code>	
<i>sqlite-memory</i> : SQLite memory used from <code>osi_get_sqlite_status</code>	
<i>sqlite-memory-highwater</i> : SQLite memory highwater since last event from <code>osi_get_sqlite_status</code>	
<i>databases</i> : number of open SQLite databases	
<i>listeners</i> : number of open TCP/IP listeners	
<i>ports</i> : number of open osi-ports	
<i>watchers</i> : number of open path watchers	
<i>cpu</i> : CPU time in seconds since last event	
<i>real</i> : elapsed time in seconds since last event	
<i>bytes</i> : Scheme heap bytes allocated since last event	
<i>gc-count</i> : number of garbage collections since last event	
<i>gc-cpu</i> : CPU time in seconds of garbage collections since last event	
<i>gc-real</i> : elapsed time in seconds of garbage collections since last event	
<i>gc-bytes</i> : Scheme heap bytes reclaimed since last event	

This event is sent every five minutes while the `statistics` gen-server is running.

Chapter 12

HTTP Interface

12.1 Introduction

The HTTP interface provides a basic implementation of the Hypertext Transfer Protocol [13]. The programming interface includes procedures for the HyperText Markup Language (HTML) version 5 [17] and JavaScript Object Notation (JSON) [3].

12.2 Theory of Operation

The HTTP interface provides a supervisor, `http-sup`, to manage the `http-listener` gen-server, the `http-cache` gen-server, and new connection processes. This structure is illustrated in Figure 12.1.

The `http-listener` is a gen-server that creates a TCP listener using `listen-tcp` and accepts new connections using `accept-tcp`. For each connection, the `http-listener` uses its supervisor to spawn and link a handler process.

Each handler reads from its input port until a CR LF occurs. Well-formed input is converted to a `<request>` tuple, and the HTTP request header and any content parameters are read.

When `Content-Length` appears in the header, the content bytes are read. If the `Content-Type` is `multipart/form-data` or `application/x-www-form-urlencoded`, the content is converted to




Figure 12.1: HTTP tree

an association list. Otherwise, parameter "unhandled-content" is included with the value of the raw bytevector of data. Each uploaded file is stored in (tmp-dir), and the association list value is #(<file> filename).

http:file-handler is then called combining the <request> query parameters and content parameters. http:file-handler logs the specific request, validates that the requested path does not include "..", retrieves a page handler from the http-cache, and invokes it. A *page handler* is a procedure which responds to a particular HTTP request. The output port is flushed after the page handler returns.

After a request is processed, all uploaded files are deleted, and the current process and connection can be reused. The system reads another request from the input port.

The http-cache is a gen-server that stores page handlers and provides a mapping from file extension to content type. It creates a directory watcher using watch-directory to invalidate the cache when anything in the (web-dir) tree changes.

The http-cache considers a path that ends in ".ss" a dynamic page loaded from (web-dir). Other paths are considered static and are sent directly over the connection using http:respond-file.

12.3 Security

The HTTP interface is written in Scheme, and therefore buffer overrun exploits cannot be used against the system.

User input should be carefully checked before calling eval or invoking a database query.

A URL which directs the system away from (web-dir) using "." could allow access to system files. http:file-handler explicitly checks for relative paths.

The HTTP interface limits incoming data to protect against large memory allocation which may crash the system. URL requests are limited to 4,096 bytes. Headers are limited to 1,048,576 bytes. Posted content is limited to 4,194,304 bytes, not including uploaded files.

The HTTP interface does not limit incoming file uploads. If the disk runs out of space, the handler process will exit with an I/O error.

12.4 Dynamic Pages

A dynamic page is a sequence of definitions followed by a sequence of expressions stored in ".ss" files in (web-dir). The definitions and expressions are placed in a **lambda** expression that is evaluated by the **interpret** system procedure. The page is responsible for sending the HTTP response. The output port is flushed after the page handler returns.

12.5 Dynamic Page Constructs

The evaluated **lambda** expression exposes the following variables to a dynamic page:

ip: binary input port
op: binary output port
request: a <request> tuple
header: an association list
params: an association list

(find-param *key*) **syntax**

Implementation: The `find-param` macro expands to `(http:find-param key params)`.

(get-param *key*) **syntax**

Implementation: The `get-param` macro expands to `(http:get-param key params)`.

(http:include "*filename*") **syntax**

The `http:include` construct includes the definitions from *filename*, a path relative to `(web-dir)` if *filename* begins with a forward slash, else relative to the directory of the current file.

Implementation: The `http:include` macro calls `read-file` and `read-bytevector` to retrieve a list of expressions that are spliced in at the same scope as the use of `http:include`. The splicing is done with `let-syntax` so that any nested `http:include` expressions are processed relative to the directory of *filename*.

12.6 Programming Interface

<request> **tuple**

method: a symbol
path: a decoded string
query: a decoded association list

http-port-number **parameter**

value: `#f` or a fixnum $0 \leq port \leq 65535$

The `http-port-number` parameter specifies whether or not `http-sup:start&link` should start the HTTP server and, if started, on what port that server should listen for connections.

(http-sup:start&link) **procedure**

returns: `#(ok pid)` | `#(error error)`

If `(http-port-number)` is not `#f`, the `http-sup:start&link` procedure creates a supervisor named `http-sup` using `supervisor:start&link` configured one-for-one with up to 10 restarts every 10 seconds. The supervisor starts the `http-cache` and `http-listener` gen-servers.

(http:get-port-number) **procedure**

returns: see below

If the `(http-port-number)` is configured to be zero, the operating system will choose an available port number. `(http:get-port-number)` uses `listener-port-number` to retrieve the actual port number that the server is listening on.

(http:find-header *name header*) **procedure**

returns: a string | #f

The `http:find-header` procedure returns the value associated with *name* in *header*. Header comparisons are case-insensitive. If *name* is not a string, exception `#(bad-arg http:find-header name)` is raised.

(http:get-header *name header*) **procedure**

returns: a string

The `http:get-header` procedure returns the value associated with *name* in *header* or exits with reason `#(invalid-header name header)`. Header comparisons are case-insensitive. If *name* is not a string, exception `#(bad-arg http:get-header name)` is raised.

(http:find-param *name params*) **procedure**

returns: a string | #f

The `http:find-param` procedure returns the value associated with *name* in *params*. Parameter comparisons are case-sensitive. If *name* is not a string, exception `#(bad-arg http:find-param name)` is raised.

(http:get-param *name params*) **procedure**

returns: a string

The `http:get-param` procedure returns the value associated with *name* in *params* or exits with reason `#(invalid-param name params)`. Parameter comparisons are case-sensitive. If *name* is not a string, exception `#(bad-arg http:get-param name)` is raised.

(http:read-header *ip limit*) **procedure**

returns: an association list

The `http:read-header` procedure reads from the binary input port *ip* until a blank line is read.

An association list is created by making a string from the characters before the first colon as the key. Non-linear white space is skipped, and the remaining characters are converted to a string value.

Reading beyond *limit* will result in exiting with reason `input-limit-exceeded`.

Failure to find a colon on any given line will result in exiting with reason `invalid-header`.

(http:read-status *ip limit*) **procedure**

returns: number | #f

The `http:read-status` procedure reads the HTTP response status line from the binary input port *op* and returns the number if well formed and `#f` otherwise. Reading beyond *limit* will result in exiting with reason `input-limit-exceeded`.

(http:write-status *op status*) **procedure**

returns: unspecified

The `http:write-status` procedure writes the HTTP response status line to the binary output port *op*.

Unless *status* is a fixnum and $100 \leq \textit{status} \leq 599$, the exception `#{bad-arg http:write-status status}` is raised.

According to HTTP [13] the status line includes a human readable reason phrase. The grammar shows that it can in fact be 0 characters long; therefore, the reason phrase is not included in this implementation.

<code>(http:write-header op header)</code>	procedure
returns: unspecified	

The `http:write-header` procedure writes the HTTP *header*, and trailing CR LF to the binary output port *op*.

header is an association list. If *header*'s keys are not strings, exception `#{bad-arg http:write-header header}` is raised.

<code>(http:respond op status header content)</code>	procedure
returns: unspecified	

The `http:respond` procedure writes the HTTP *status* and *header* to binary output port *op* using `http:write-status` and `http:write-header`, adding `Content-Length` to the *header*. When `Cache-Control` is not present in *header*, it is added with value `no-cache`. The *content* is then written, and the output port is flushed.

content is a bytevector.

<code>(http:respond-file op status header filename)</code>	procedure
returns: unspecified	

The `http:respond-file` procedure writes the HTTP *status* and *header* to binary output port *op* using `http:write-status` and `http:write-header`, adding `Content-Length` to *header*. The `Cache-Control` header is added, if it is not already present, with value `max-age=3600`. The `Content-Type` header is added if it is not already present and the extension of *filename* matches (case insensitively) an extension in the `mime-types` file of `(web-dir)`. Each line of `mime-types` has the form `("extension" . "Content-Type")`. The content of the file is streamed to the output port so that the file does not need to be loaded into memory. The output port is flushed.

<code>(http:percent-encode s)</code>	procedure
returns: an encoded string	

The `http:percent-encode` procedure writes the characters A–Z, a–z, 0–9, hyphen, underscore, period, and `~`. Other characters are converted to a `%` prefix and two digit hexadecimal representation.

<code>(html:encode s)</code>	procedure
<code>(html:encode op s)</code>	

returns: see below

The `html:encode` procedure converts special character entities in string *s*.

input	output
"	";
&	&
<	<
>	>

The single argument form of `html:encode` returns an encoded string.

The two argument form of `html:encode` sends the encoded string to the textual output port *op*.

<pre>(html->string x)</pre>	procedure
<pre>(html->string op x)</pre>	
returns: see below	

The `html->string` procedure transforms an object into HTML. The transformation, *H*, is described below:

<i>x</i>	<i>H(x)</i>
()	nothing
#!void	nothing
<i>string</i>	<i>E(string)</i>
<i>number</i>	<i>number</i>
(begin <i>pattern</i> ...)	<i>H(pattern)</i> ...
(cdata <i>string</i> ...)	[!CDATA[<i>string</i> ...]]
(html5 [(@ <i>attr</i> ...)] <i>pattern</i> ...)	<!DOCTYPE html><html <i>A(attr)</i> ...> <i>H(pattern)</i> ...</html>
(raw <i>string</i> ...)	<i>string</i> ...
(script [(@ <i>attr</i> ...)] <i>string</i> ...)	<script <i>A(attr)</i> ...> <i>string</i> ...</script>
(style [(@ <i>attr</i> ...)] <i>string</i> ...)	<style <i>A(attr)</i> ...> <i>string</i> ...</style>
(tag [(@ <i>attr</i> ...)] <i>pattern</i> ...)	<tag <i>A(attr)</i> ...> <i>H(pattern)</i> ...</tag>
(void-tag [(@ <i>attr</i> ...)])	<void-tag <i>A(attr)</i> ...>

E denotes the `html:encode` function.

A *void-tag* is one of `area`, `base`, `br`, `col`, `embed`, `hr`, `img`, `input`, `keygen`, `link`, `menuitem`, `meta`, `param`, `source`, `track`, or `wbr`. A *tag* is any other symbol.

The attribute transformation, *A*, is described below, where *key* is a symbol:

<i>attr</i>	<i>A(attr)</i>
#!void	nothing
(<i>key</i>)	<i>key</i>
(<i>key string</i>)	<i>key</i> ="E(<i>string</i>)"
(<i>key number</i>)	<i>key</i> ="number"

The single argument form of `html->string` returns an encoded HTML string.

The two argument form of `html->string` sends the encoded HTML string to the textual output port *op*.

Input that does not match the specification causes a `#(bad-arg html->string x)` exception to be raised.

<code>(html->bytevector <i>x</i>)</code>	procedure
returns: a bytevector	

The `html->bytevector` procedure calls `html->string` on *x* using a bytevector output port transcoded using `(make-utf8-transcoder)` and returns the resulting bytevector.

12.6.1 JavaScript Object Notation

This implementation translates JavaScript types into the following Scheme types:

JavaScript	Scheme
<code>true</code>	<code>#t</code>
<code>false</code>	<code>#f</code>
<code>null</code>	<code>#\nul</code>
<i>string</i>	<i>string</i>
<i>number</i>	<i>number</i>
<i>array</i>	<i>list</i>
<i>object</i>	hashtable mapping case-sensitive strings to values

This implementation does not range check values to ensure that a JavaScript implementation can interpret the data.

<code>(json:extend-object <i>ht</i> [<i>key value</i>] ...)</code>	syntax
--	---------------

The `json:extend-object` construct adds the *key* / *value* pairs to the hashtable *ht* using `hashtable-set!`. The resulting expression returns *ht*.

<code>(json:make-object [<i>key value</i>] ...)</code>	syntax
--	---------------

The `json:make-object` construct expands into a call to `json:extend-object` with a new hashtable.

<code>(json:read <i>ip</i> [<i>custom-inflate</i>])</code>	procedure
returns: a Scheme object	

The `json:read` procedure reads characters from the textual input port *ip* and returns an appropriate Scheme object. When `json:read` encounters a JSON object, it builds the corresponding hashtable and calls *custom-inflate* to perform application-specific conversion. By default, *custom-inflate* is the identity function.

The following exceptions may be raised:

- `invalid-surrogate-pair`
- `unexpected-eof`
- `#(unexpected-input data input-position)`

<code>(json:write <i>op</i> <i>x</i> [<i>custom-write</i>])</code>	procedure
--	------------------

returns: unspecified

The `json:write` procedure writes the object *x* to the textual output port *op* in JSON format. JSON objects are sorted by key using `string<?` to provide stable output. Scheme fixnums, bignums, and finite flonums may be used as numbers.

The optional *custom-write* procedure may intervene to handle lists and hashtables differently or to handle objects that have no direct JSON counterpart. If *custom-write* does not handle a given object, it should return false to let `json:write` proceed normally. The *custom-write* procedure is called with three arguments: the textual output port *op*, the Scheme object *x*, and a writer procedure *wr* that should be used to write the values of arbitrary Scheme objects. The *wr* procedure is equivalent to `(lambda (op x) (json:write op x custom-write))`.

If an object cannot be formatted, `#(invalid-datum x)` is raised.

`(json:write-object op wr [key value] ...)` **syntax**

Given a textual output port *op* and a writer procedure *wr*, the `json:write-object` construct writes a JSON object with the given *key* / *value* pairs to *op*, sorted by key using `string<?`. Each *key* must be a distinct string literal. The *wr* procedure takes *op* and an object *x*, just like the *wr* procedure that is passed to `json:write`'s *custom-write* procedure.

The following are equivalent, provided the keys are string literals.

```
(json:write op (json:make-object [key value] ...))  
(json:write-object op json:write [key value] ...)
```

The latter trades code size and compile time for run-time efficiency. At compile time, `json:write-object` sorts the keys and preformats the strings that will separate values. For example:

```
(json:write-object op wr ["bar" e0] ["foo" e1] ["ace" e2])
```

expands to:

```
(let ([op op] [wr wr])  
  (display "{\"ace\":\"" op)  
  (wr op e2)  
  (display "\",\"bar\":\"" op)  
  (wr op e0)  
  (display "\",\"foo\":\"" op)  
  (wr op e1)  
  (write-char #\} op))
```

`(json:object->bytevector x [custom-write])` **procedure**

returns: a bytevector

The `json:object->bytevector` procedure calls `json:write` on *x* with the optional *custom-write*, if any, using a bytevector output port transcoded using `(make-utf8-transcoder)` and returns the resulting bytevector.

`(json:object->string x [custom-write])` **procedure**

returns: a JSON formatted string

The `json:object->string` procedure creates a string output port, calls `json:write` on *x* with the optional *custom-write*, if any, and returns the resulting string.

`(json:string->object x [custom-inflate])` **procedure**

returns: a Scheme object

The `json:string->object` procedure creates a string input port on *x*, calls `json:read` with the optional *custom-inflate*, if any, and returns the resulting Scheme object after making sure the rest of the string is only whitespace.

12.7 Published Events

`<http-request>` **event**

timestamp: timestamp from `erlang:now`

pid: handler process

host: the IP address of the client

method: `<request>` method

path: `<request>` path

header: an association list

params: an association list

Chapter 13

Command Line Interface

13.1 Introduction

The command-line interface (`cli`) provides parsing of command-line arguments as well as consistent usage of common options and display of help.

13.2 Theory of Operation

Many programs parse command-line arguments and perform actions based on them. The `cli` library helps to make programs that process arguments and display help simple and consistent. Command-line arguments are parsed left to right in a single pass. Command-line interface specifications, or `cli-specs`, are used for parsing and error checking a command line, displaying one-line usage, and displaying a full help summary.

Arguments may be preceded by a single dash (`-`), a double dash (`--`), or no dash at all. A single dash precedes short, single character arguments. The API does not allow numbers as they could be mistaken as a negative numerical value supplied to another argument. A double dash precedes longer, more descriptive arguments, `--repl` for example. Positional arguments are not preceded by any dashes. As arguments with dashes are consumed, the remaining arguments are matched against the positional specifications in order.

Argument specifications include a type such as: `bool`, `count`, `string`, and `list`. A set of `bool` and `count` arguments can be specified together (`-abc` is equivalent to `-a -b -c`). Arguments of type `list` collect values in left to right order.

The API does not directly support sub-commands and alternate usage help text. These can be implemented using the primitives provided. The implementations of `swish-build` and `swish-test` provide examples of advanced command-line handling.

In the following REPL transcript, we define `example-cli` using `cli-specs`. We then set the `command-line-arguments` parameter as they would be for an application. Calling `parse-command-line-arguments` returns a procedure, `opt`, which we can use to access the parsed command-line values. Finally, we use `display-help` to display the automatically generated help.

```

> (define example-cli
  (cli-specs
    default-help
    ["verbose" -v count "indicates verbosity level"]
    ["output" -o (string "<output>") "print output to an <output> file"]
    ["repl" --repl bool "start a repl"]
    ["files" (list "<file>" ...) "a list of input files"])))
> (command-line-arguments '("-vvv" "-o" "file.out" "file.in"))
> (define opt (parse-command-line-arguments example-cli))
> (opt "verbose")
3
> (opt "output")
"file.out"
> (opt "files")
("file.in")
> (display-help "sample" example-cli)
Usage: sample [-hv] [-o <output>] [--repl] <file> ...

-h, --help      display this help and exit
-v              indicates verbosity level
-o <output>     print output to an <output> file
--repl          start a repl
<file> ...      a list of input files

```

Putting the parts together into `sample.ss`, we have a working example albeit incomplete.

```
#!/usr/bin/env swish
```

```

(define example-cli
  (cli-specs
    default-help
    ["verbose" -v count "indicates verbosity level"]
    ["output" -o (string "<output>") "print output to an <output> file"]
    ["repl" --repl bool "start a repl"]
    ["files" (list "<file>" ...) "a list of input files"])))

(let ([opt (parse-command-line-arguments example-cli)])
  (when (opt "help")
    (display-help "sample" example-cli)
    (exit 0))
  (let ([verbosity (or (opt "verbose") 0)])
    (when (> verbosity 0)
      (printf "showing verbosity level: ~a~n" verbosity)))
  (when (opt "repl")
    (new-cafe)))

```

13.3 Programming Interface

```
(cli-specs                                     syntax
 [default-help]
 (name [short] [long] type help
  [(conflicts conflicts)]
  [(requires requires)]
  [(usage [visibility] [how])])
 ...)
```

expands to:

a list of `<arg-spec>` tuples

The `cli-specs` macro simplifies the creation of the `<arg-spec>` tuples. The `<arg-spec>` `name` field uniquely identifies a specification, and is used to retrieve parsed argument values and check constraints.

name: a string to identify the argument

short: a symbol of the form `-x`, where *x* is a single character, see below

long: a symbol of the form `--x`, where *x* is a string

type: see Figure 13.1

help: a string or list of strings that describes the argument

conflicts: a list of `<arg-spec>` names

requires: a list of `<arg-spec>` names

To specify `-i` or `-I` for *short*, use `|-i|` and `|-I|` respectively to prevent Chez Scheme from reading them as the complex number $0 - 1i$.

Type	Result
<code>bool</code>	<code>#t</code>
<code>count</code>	a positive integer
<code>(string x)</code>	a string
<code>(list x)</code>	a list of one item
<code>(list x ...)</code>	a list of one or more items up to the next argument
<code>(list . x)</code>	a list of the rest of the arguments

For each type where *x* is specified, *x* is a string that is used in the help display.

Figure 13.1: Command-line argument types

The `list` types can support multiple *x* arguments, for instance `(list "i1" "i2")` would specify a list of two arguments.

```
visibility → show
           | hide
           | fit
```

When printing the help usage line, a *visibility* of `show` means the argument must be displayed. `hide` forces the argument to be hidden. `fit` displays the argument if it fits on the line.

```

how →  short
      |  long
      |  opt
      |  req

```

The *how* expands into input of the `format-spec` procedure according to Figure 13.2.

Keyword	Expands into:
short	(opt (and short args))
long	(opt (and long args))
opt	(opt (and (or short long) args))
req	(req (and (or short long) args))

Figure 13.2: `cli-specs` *how* field

For options with *short* or *long* specified, `fit` and `opt` are the defaults. For other options, `show` and `req` are the defaults.

conflicts is a list of specification names that prevent this argument from processing correctly. When multiple command-line arguments are specified that are in conflict, an exception is raised.

requires is a list of other specification names that are necessary for this argument to be processed correctly. Unless all the required command-line arguments are specified, an exception is raised.

The `conflicts`, `requires`, and `usage` clauses may be specified in any order.

```

(display-help exe-name specs [args] [op]) procedure
returns: unspecified

```

The `display-help` procedure is equivalent to calling `display-usage` with a prefix of "Usage:" followed by `display-options`.

```

(display-options specs [args] [op]) procedure
returns: unspecified

```

For each specification in *specs*, the `display-options` procedure renders two columns of output to *op*, which defaults to the current output port. The first column renders the short and long form of the argument with its additional inputs. The second column renders the `<arg-spec>` `help` field and will automatically wrap if the `help-wrap-width` is exceeded.

If an *args* hash table is specified, the specified value appended to the second column. This is useful for displaying default or current values.

```

(display-usage prefix exe-name specs [width] [op]) procedure
returns: unspecified

```

The `display-usage` procedure displays the first line of help output to *op*, which defaults to the current output port. It starts with *prefix* and *exe-name* then attempts to fit *specs* onto the line using `format-spec`. When the line will exceed *width* characters, some arguments may collapse to `[options]`.

A *width* of `#f` defaults the line width to `help-wrap-width`.

(format-spec spec [how])	procedure
returns: a string	

The **format-spec** procedure is responsible rendering *spec* as a string as specified by *how*. **format-spec** can display dashes in front of arguments, ellipses on list types, and brackets around optional arguments. A *how* of **#f** defaults to the **<arg-spec> usage** field.

<i>how</i>	Return value:														
short	"-x" if <i>spec short</i> is x, else #f														
long	"--x" if <i>spec long</i> is x, else #f														
args	The <i>spec type</i> is evaluated as follows: <table border="0" style="margin-left: 20px;"> <tr> <th><i>type</i></th><th>Return value:</th></tr> <tr> <td>bool</td><td>#f</td></tr> <tr> <td>count</td><td>#f</td></tr> <tr> <td>(string x)</td><td>"x"</td></tr> <tr> <td>(list x)</td><td>"x"</td></tr> <tr> <td>(list x ...)</td><td>"x ..."</td></tr> <tr> <td>(list . x)</td><td>"x ..."</td></tr> </table>	<i>type</i>	Return value:	bool	#f	count	#f	(string x)	"x"	(list x)	"x"	(list x ...)	"x ..."	(list . x)	"x ..."
<i>type</i>	Return value:														
bool	#f														
count	#f														
(string x)	"x"														
(list x)	"x"														
(list x ...)	"x ..."														
(list . x)	"x ..."														
(or how ...)	Recur and use the first non- #f														
(and how ...)	Recur and concatenate all non- #f values														
(opt how)	Recur and surround the result with square brackets [] if non- #f														
(req how)	Recur and use the result														

help-wrap-width	parameter
value: a positive fixnum	

The **help-wrap-width** parameter specifies the default width for **display-usage** and **display-options**.

(parse-command-line-arguments specs [ls] [fail])	procedure
returns: a procedure	

The **parse-command-line-arguments** procedure processes the elements of *ls* from left to right in a single pass. As it scans each *x* in *ls*, the parser must find a suitable *s* within *specs*. If a suitable *s* cannot be found, the parser reports an error by calling *fail*. Based on the type of *s*, the parser may consume additional elements following *x*. The type of *s* determines what data the parser records for that argument. When *s* is satisfied, the parser continues scanning the remaining elements of *ls*.

The parser returns a procedure *p* that accepts zero or one argument. When called with no arguments, *p* returns a hash table that maps the name of each *s* found while processing *ls* to the data recorded for that argument. When called with the name of an element *s* in *specs*, *p* returns the data, if any, recorded for that name in the hash table or else **#f**. If a particular *s* was not found while processing *ls*, the internal hash table has no entry for the name of *s* and *p* returns **#f** when given that name. If called with a name that is not in *specs*, *p* raises an exception.

The following table summarizes the parser's behavior.

<code><arg-spec></code> type	extra arguments consumed / recorded	return value of (<i>p name</i>)
<code>bool</code>	none	#t
<code>count</code>	none	an exact positive integer
<code>(string x)</code>	one	a string
<code>(list x₀ ... x_n ...)</code>	<i>n</i> or more, up to the next option	a list of strings
<code>(list x₀ ... x_n . rest)</code>	at least <i>n</i> and all remaining	a list of strings

By default *ls* is the value of `(command-line-arguments)` and *fail* is a procedure that applies **errorf** to its arguments. Providing a *fail* procedure allows a developer to accumulate parsing errors without necessarily generating exceptions.

<arg-spec>	tuple
-------------------------	--------------

name: a string to use as the key of the output hash table
type: see Figure 13.1
short: **#f** | a character
long: **#f** | a string
help: a string or list of strings describing argument
conflicts: a list of `<arg-spec>` names
requires: a list of `<arg-spec>` names
usage: a list containing one *visibility* symbol and a `format-spec` *how* expression

Bibliography

- [1] Joe Armstrong. *Making reliable distributed systems in the presence of software errors*. PhD thesis, The Royal Institute of Technology, Stockholm, Sweden, 2003.
- [2] Joe Armstrong. *Programming Erlang—Software for a Concurrent World*. The Pragmatic Bookshelf, 2007.
- [3] T. Bray. The JavaScript Object Notation (JSON) Data Interchange Format. RFC 7159 (Proposed Standard), March 2014. <http://www.ietf.org/rfc/rfc7159.txt>.
- [4] Carl Bruggeman, Oscar Waddell, and R. Kent Dybvig. Representing Control in the Presence of One-Shot Continuations. In *Proceedings of the ACM SIGPLAN '96 Conference on Programming Language Design and Implementation*, pages 99–107, 1996.
- [5] C99 — Wikipedia, The Free Encyclopedia. <https://en.wikipedia.org/w/index.php?title=C99&oldid=813613099>.
- [6] R. Kent Dybvig. *Chez Scheme Version 9 User's Guide*. Cadence Research Systems, 2017. <https://cisco.github.io/ChezScheme/csug9.5/csug.html>.
- [7] R. Kent Dybvig, Robert Hieb, and Carl Bruggeman. Syntactic Abstraction in Scheme. *Lisp and Symbolic Computation*, 5:83–110, 1992.
- [8] Erlang. <http://www.erlang.org/>.
- [9] Erlang dict module. <http://www.erlang.org/doc/man/dict.html>.
- [10] Erlang gen_server module. http://www.erlang.org/doc/man/gen_server.html.
- [11] Erlang queue module. <http://www.erlang.org/doc/man/queue.html>.
- [12] Erlang supervisor module. <http://www.erlang.org/doc/man/supervisor.html>.
- [13] R. Fielding and J. Reschke. Hypertext Transfer Protocol (HTTP/1.1): Message Syntax and Routing. RFC 7230 (Proposed Standard), June 2014. <http://www.ietf.org/rfc/rfc7230.txt>.
- [14] William G. Griswold and Gregg M. Townsend. The design and implementation of dynamic hashing for sets and tables in icon. *Software Practice and Experience*, 23(4):351–367, April 1993.
- [15] Robert Hanmer. *Patterns for Fault Tolerant Software*. Wiley Publishing, 2007.

- [16] C. Hewitt, P. Bishop, and R. Steiger. A Universal Modular ACTOR Formalism for Artificial Intelligence. In *Proceedings of the Third IJCAI*, pages 235–245, Stanford, MA, 1973.
- [17] Ian Hickson. HTML 5, October 2014. <http://www.w3.org/TR/2014/REC-html5-20141028/>.
- [18] Robert Hieb, R. Kent Dybvig, and Carl Bruggeman. Representing Control in the Presence of First-Class Continuations. In *Proceedings of the ACM SIGPLAN '90 Conference on Programming Language Design and Implementation*, pages 66–77, 1990.
- [19] libuv. <http://libuv.org/>.
- [20] Michael Owens. *The Definitive Guide to SQLite*. Apress, 2006.
- [21] E. Resnick. Internet Message Format, April 2001. <http://www.ietf.org/rfc/rfc2822.txt>.
- [22] Dorai Sitaram. pregexp: Portable Regular Expressions for Scheme and Common Lisp, 2005. <http://ds26gte.github.io/pregexp/index.html>.
- [23] SQLite. <http://www.sqlite.org/>.

List of Figures

1.1	Supervision Tree	7
3.1	Pattern Grammar	18
5.1	Event flow	43
12.1	HTTP tree	73
13.1	Command-line argument types	84
13.2	<code>cli-specs</code> <i>how</i> field	85

Index

<arg-spec>, 87
<child-end>, 55
<child-start>, 55
<event-logger>, 68
<gen-server-debug>, 38
<gen-server-terminating>, 38
<handler>, 43
<http-request>, 81
<request>, 75
<statistics>, 72
<supervisor-error>, 55
<system-attributes>, 70
<transaction-retry>, 60

absolute-path, 25
add-finalizer, 20
application, 57
 exit-process, 58
 handle-call, 57
 handle-cast, 57
 handle-info, 57
 init, 57
 state, 57
 terminate, 57
application:shutdown, 58
application:start, 58

bad-arg, 20
binary->utf8, 23

catch, 16
cli-specs, 84
close-directory-watcher, 29
close-osi-port, 26
close-tcp-listener, 30
coerce, 70
columns, 64
connect-tcp, 30
connect-usb, 24
console-event-handler, 20

create-client-pipe, 24
create-directory-path, 25
create-file, 26
create-file-port, 25
create-server-pipe, 24
create-table, 69
create-watched-process, 25
current-exit-reason->english, 32

db
 handle-call, 61
 handle-cast, 61
 handle-info, 61
 init, 61
 parameters, 60
 state, 60
 terminate, 61
db:filename, 62
db:log, 62
db:start&link, 62
db:stop, 62
db:transaction, 62
dbg, 20
define-simple-events, 69
define-state-tuple, 37
define-tuple, 22
demonitor, 17
demonitor&flush, 17
directory watcher, 12, 15, 29
directory-watcher, 12
directory-watcher guardian, 15, 29
directory-watcher-path, 29
display-help, 85
display-options, 85
display-usage, 85

ends-with-ci?, 33
ends-with?, 33
erlang:now, 21
error pair, 8

- event-loop, 13
- event-mgr
 - handle-call, 43
 - handle-info, 44
 - init, 43
 - state, 42
 - terminate, 43
- event-mgr:add-handler, 44
- event-mgr:flush-buffer, 45
- event-mgr:notify, 45
- event-mgr:set-log-handler, 45
- event-mgr:start&link, 44
- execute, 63
- execute-sql, 63
- exit-reason->english, 32
- finalizer, 13, 20
- find-files, 27
- find-param, 75
- force-close-output-port, 23
- format-rfc2822, 33
- format-spec, 85
- gatekeeper, 13, 46
 - <mutex>, 46
 - deadlock?, 47
 - enter-mutex, 47
 - handle-call, 47
 - handle-cast, 47
 - handle-info, 47
 - init, 46
 - leave-mutex, 47
 - state, 46
 - terminate, 47
- gatekeeper:enter, 48
- gatekeeper:leave, 48
- gatekeeper:start&link, 48
- gen-server:call, 36
- gen-server:cast, 37
- gen-server:debug, 37
- gen-server:reply, 37
- gen-server:start, 36
- gen-server:start&link, 36
- get-datum/annotations-all, 28
- get-file-size, 26
- get-param, 75
- get-registered, 16
- get-source-offset, 28
- handle-call, 39
- handle-cast, 40
- handle-info, 40
- help-wrap-width, 86
- hook-console-input, 27
- ht:delete, 31
- ht:fold, 31
- ht:is?, 32
- ht:keys, 32
- ht:make, 32
- ht:ref, 32
- ht:set, 32
- ht:size, 32
- html->bytevector, 78
- html->string, 78
- html:encode, 77
- http-port-number, 75
- http-sup:start&link, 75
- http:find-header, 75
- http:find-param, 76
- http:get-header, 76
- http:get-param, 76
- http:get-port-number, 75
- http:include, 75
- http:percent-encode, 77
- http:read-header, 76
- http:read-status, 76
- http:respond, 77
- http:respond-file, 77
- http:write-header, 77
- http:write-status, 76
- init, 39
- io-error, 24
- join, 33
- json:extend-object, 79
- json:make-object, 79
- json:object->bytevector, 80
- json:object->string, 80
- json:read, 79
- json:string->object, 81
- json:write, 79
- json:write-object, 80
- kill, 16
- lazy-execute, 63
- link, 16

- listen-tcp, 29
- listener, 13
- listener guardian, 15, 30
- listener-address, 30
- listener-port-number, 30
- log-db:get-instance-id, 69
- log-db:setup, 68
- log-db:start&link, 68
- log-db:version, 69

- make-process-parameter, 21
- make-utf8-transcoder, 23
- match, 19
- match-let*, 19
- mon, 12
- monitor, 17
- monitor?, 17
- move-file, 27
- msg, 12

- on-exit, 21
- open-file-to-append, 27
- open-file-to-read, 27
- open-file-to-replace, 27
- open-file-to-write, 27
- open-utf8-bytevector, 28
- osi-port, 12
- osi-port guardian, 15, 24, 25, 30
- osi_get_argv, 9
- osi_get_bytes_used, 9
- osi_get_callbacks, 9
- osi_get_error_text, 9
- osi_get_hostname, 9
- osi_get_hrttime, 9
- osi_get_time, 9
- osi_is_tick_over, 10
- osi_list_uv_handles, 10
- osi_make_uuid, 10
- osi_set_argv, 10
- osi_set_tick, 10

- parse-command-line-arguments, 86
- parse-sql, 64
- path-combine, 28
- pcb, 12
- pps, 20
- process-id, 20
- process-trap-exit, 20
- process?, 19

- profile-me, 21

- q, 11
- queue:add, 31
- queue:add-front, 31
- queue:drop, 31
- queue:empty, 31
- queue:empty?, 31
- queue:get, 31

- read-bytevector, 28
- read-file, 28
- read-osi-port, 25
- receive, 18
- register, 15
- registrar, 13
- run queue, 14, 18

- scheduler, 13
- self, 19
- send, 18
- sleep queue, 14
- spawn, 15
- spawn&link, 15
- split, 33
- split-n, 33
- sqlite:bind, 64
- sqlite:close, 64
- sqlite:columns, 65
- sqlite:execute, 65
- sqlite:expanded-sql, 65
- sqlite:finalize, 65
- sqlite:open, 65
- sqlite:prepare, 65
- sqlite:sql, 65
- sqlite:step, 65
- starts-with-ci?, 34
- starts-with?, 34
- statistics:resume, 71
- statistics:start&link, 71
- statistics:suspend, 71
- string->uuid, 10
- supervisor
 - handle-call, 50
 - handle-cast, 44, 51
 - handle-info, 52
 - init, 50
 - state, 50
 - terminate, 50

- supervisor:delete-child, 54
- supervisor:get-children, 54
- supervisor:restart-child, 54
- supervisor:start-child, 53
- supervisor:start&link, 53
- supervisor:terminate-child, 54
- swish-event-logger, 69
- swish-exit-reason->english, 33
- symbol-append, 34

TCP listener, 13, 15, 29, 30

terminate, 41

transaction, 63

tuple, 22

- copy, 23

- copy*, 23

- define-tuple, 22

- field accessor, 22

- is?, 23

- make, 22

- open, 22

unlink, 17

unregister, 15

uuid->string, 10

watch-directory, 29

watcher:shutdown-children, 55

watcher:start-child, 55

watcher:start&link, 55

whereis, 15

with-db, 62

with-gatekeeper-mutex, 48

with-sfd-source-offset, 28

write-osi-port, 26